

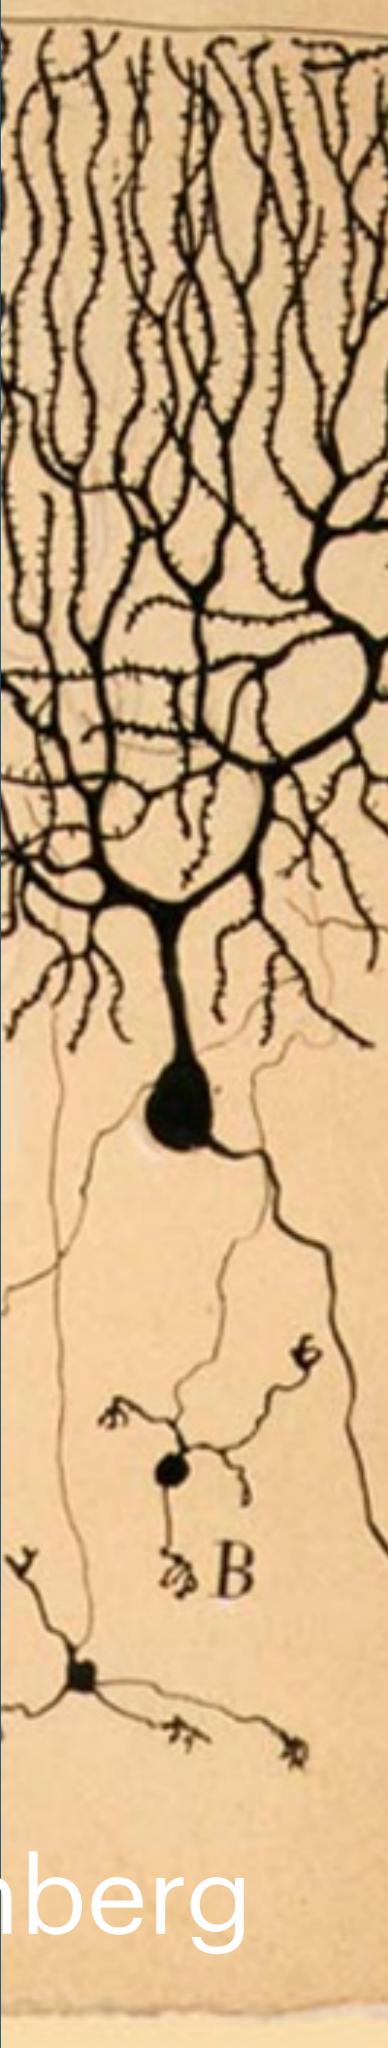
# Machine Learning in Three Easy Lessons

...with examples related  
to physics

Florian Marquardt

Max Planck Institute for the Science of Light  
Friedrich-Alexander-Universität Erlangen-Nürnberg

(Pictures:Wikimedia Commons)

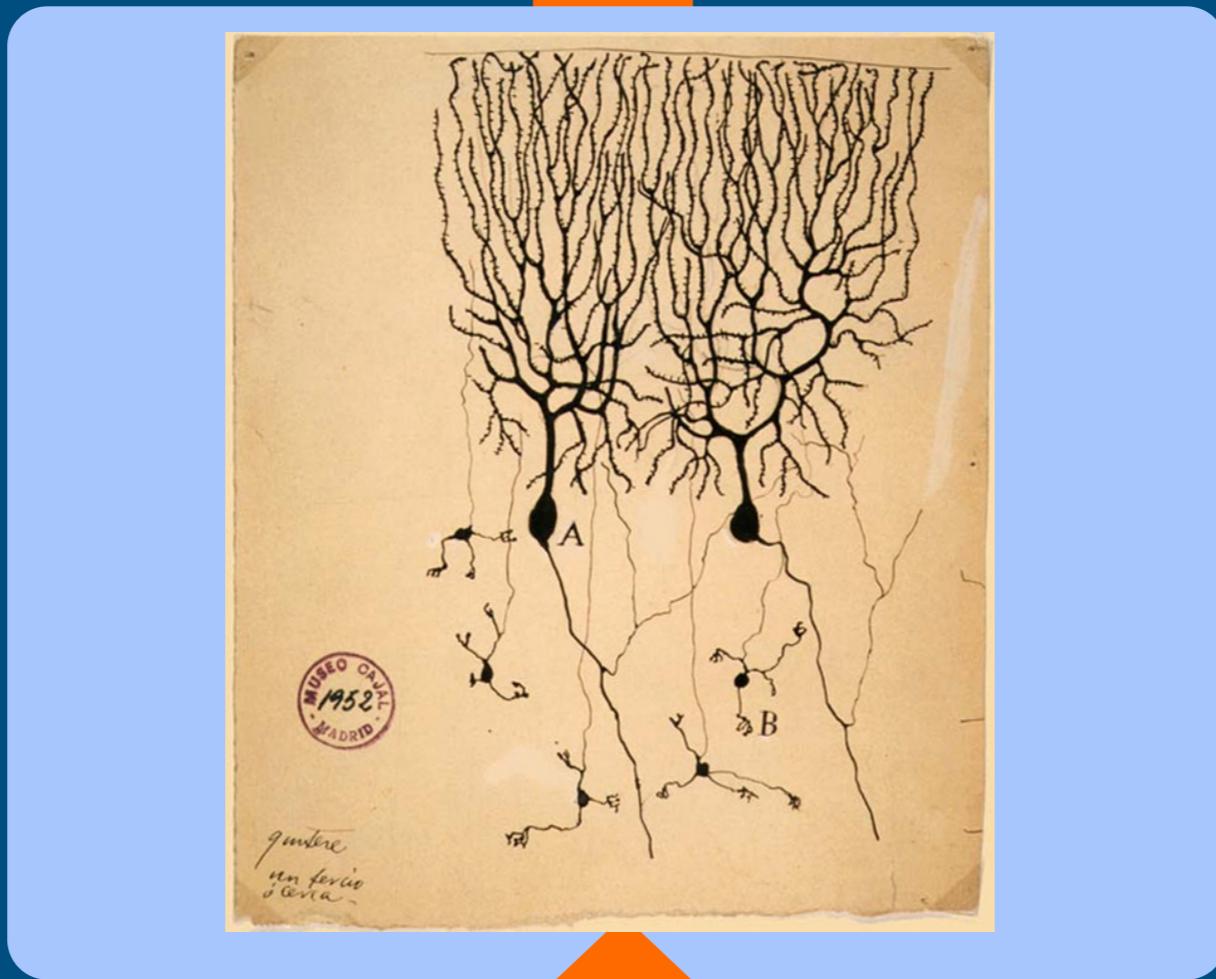
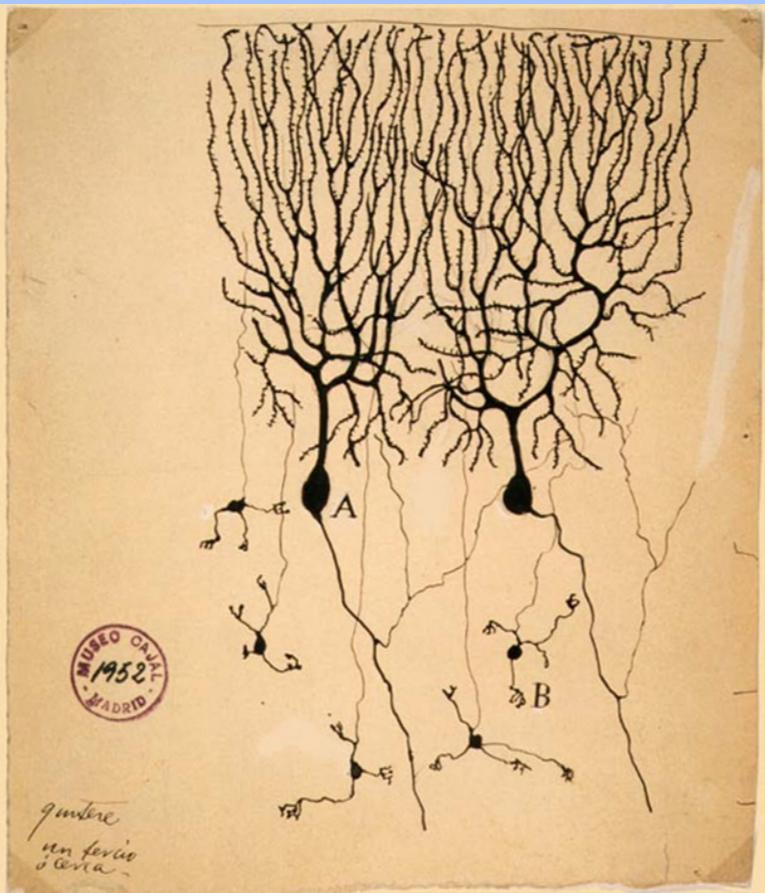




(Picture:Wikimedia Commons/Unipro)

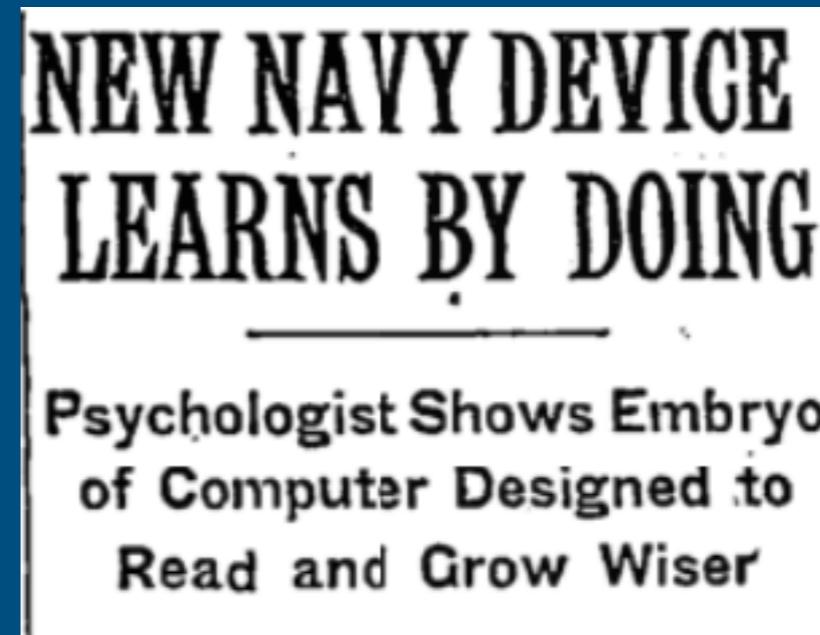
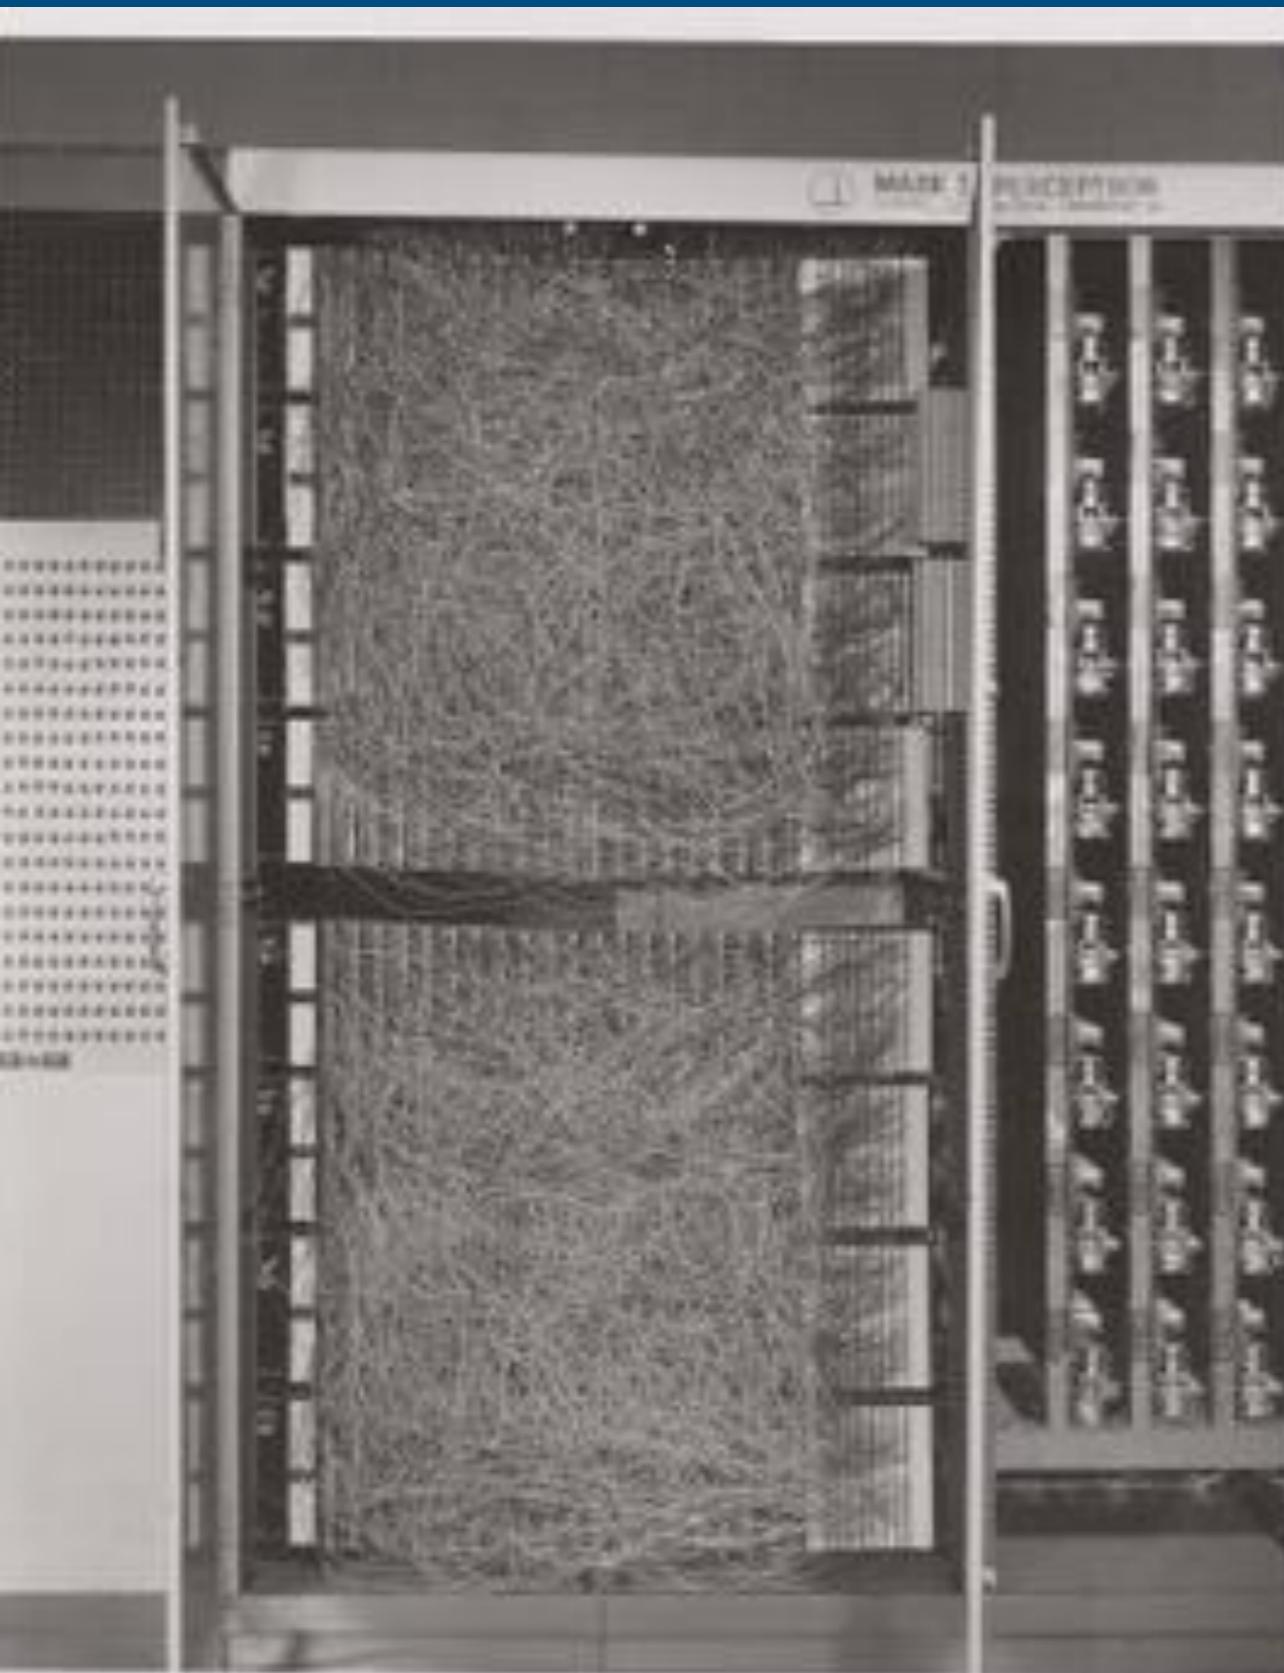
# Neural Networks

# “light bulb”



(Pictures:Wikimedia Commons)

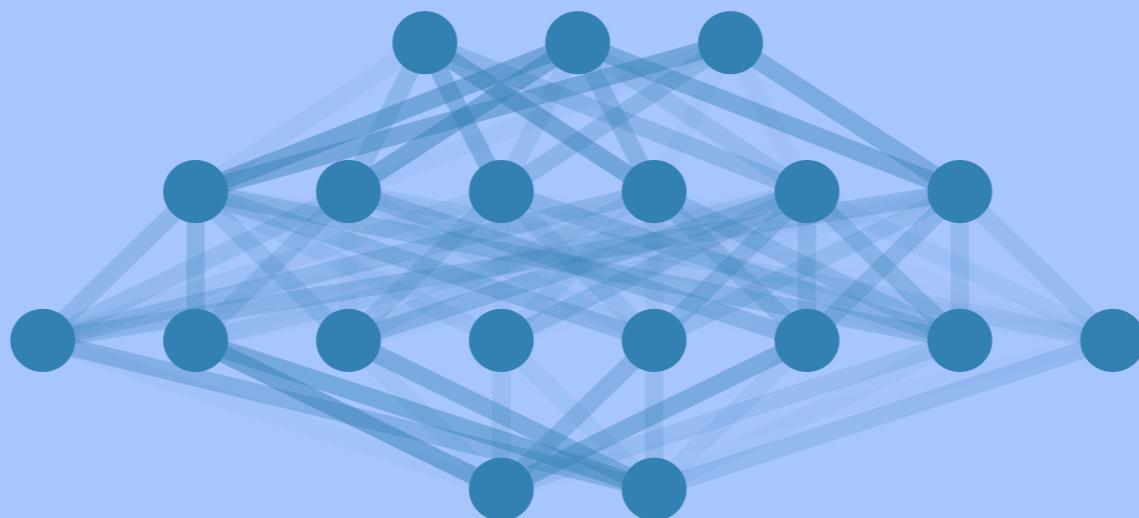
# The Perceptron Rosenblatt 1958



"Later Perceptrons will be able to recognize people and call out their names and instantly translate speech in one language to speech or writing in another language, it was predicted."  
(New York Times 1958)

(Pictures:Wikimedia Commons)

“light bulb”



Artificial Neural Network



(Pictures:Wikimedia Commons)

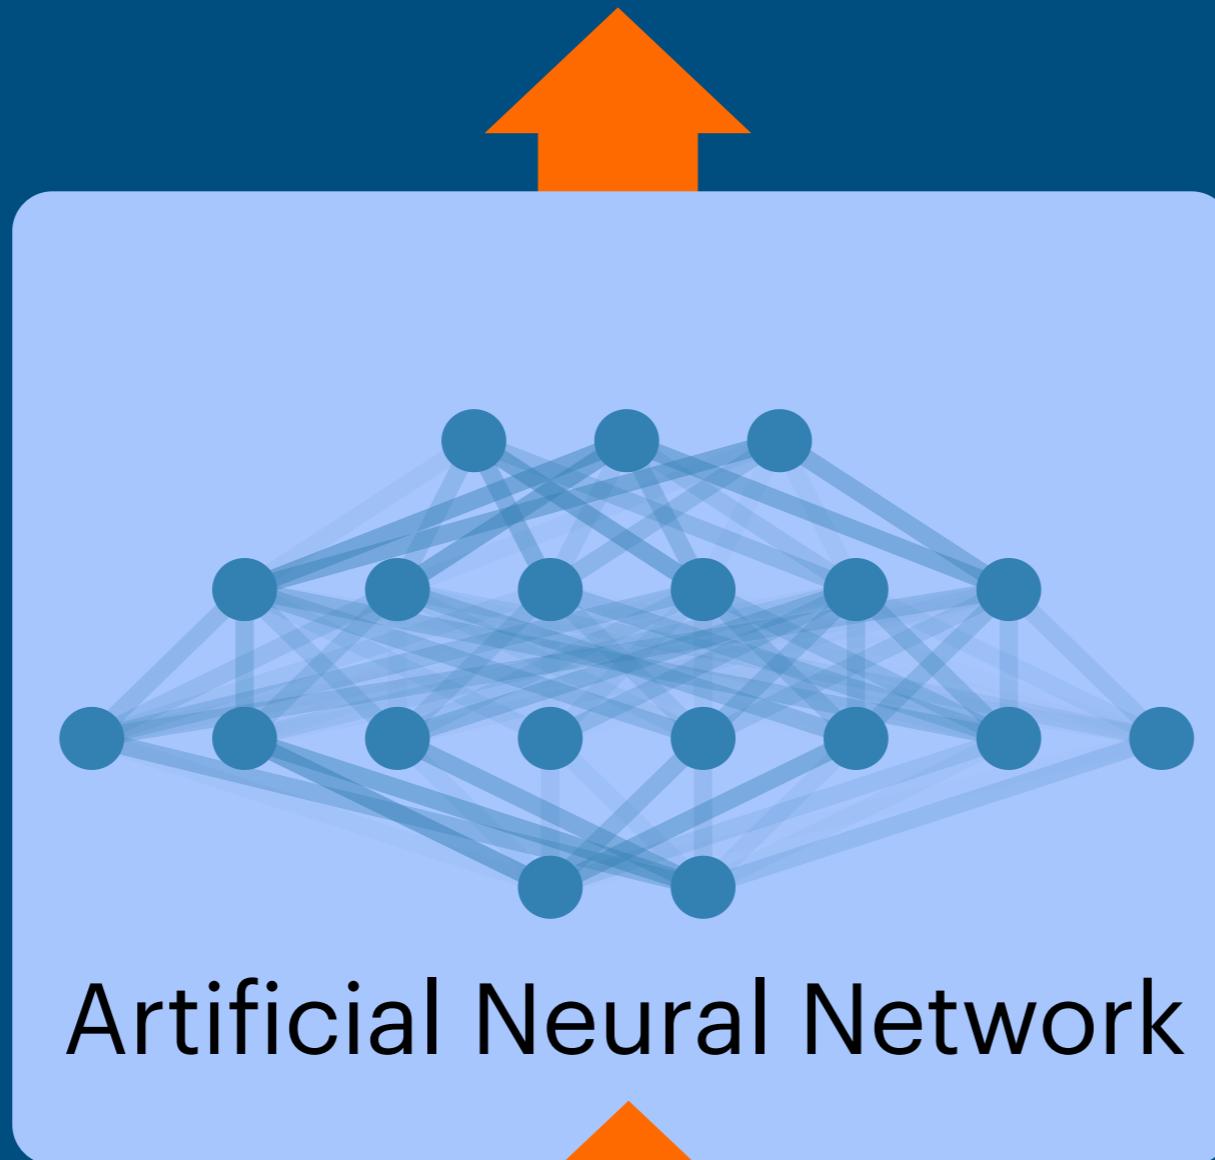
# 2012: A deep neural network beats all other approaches in image recognition



Pictures: [image-net.org](http://image-net.org)

...by now: ChatGPT, DALL-E, AlphaFold, ...

"physical parameters: damping=0.5, nonlinearity=0.1,..."



experimental  
measurements



# These lectures: Machine Learning in Three Easy Lessons

close to physics

with hands-on tutorials (no installation required)  
prior knowledge: python and numpy

# My goal

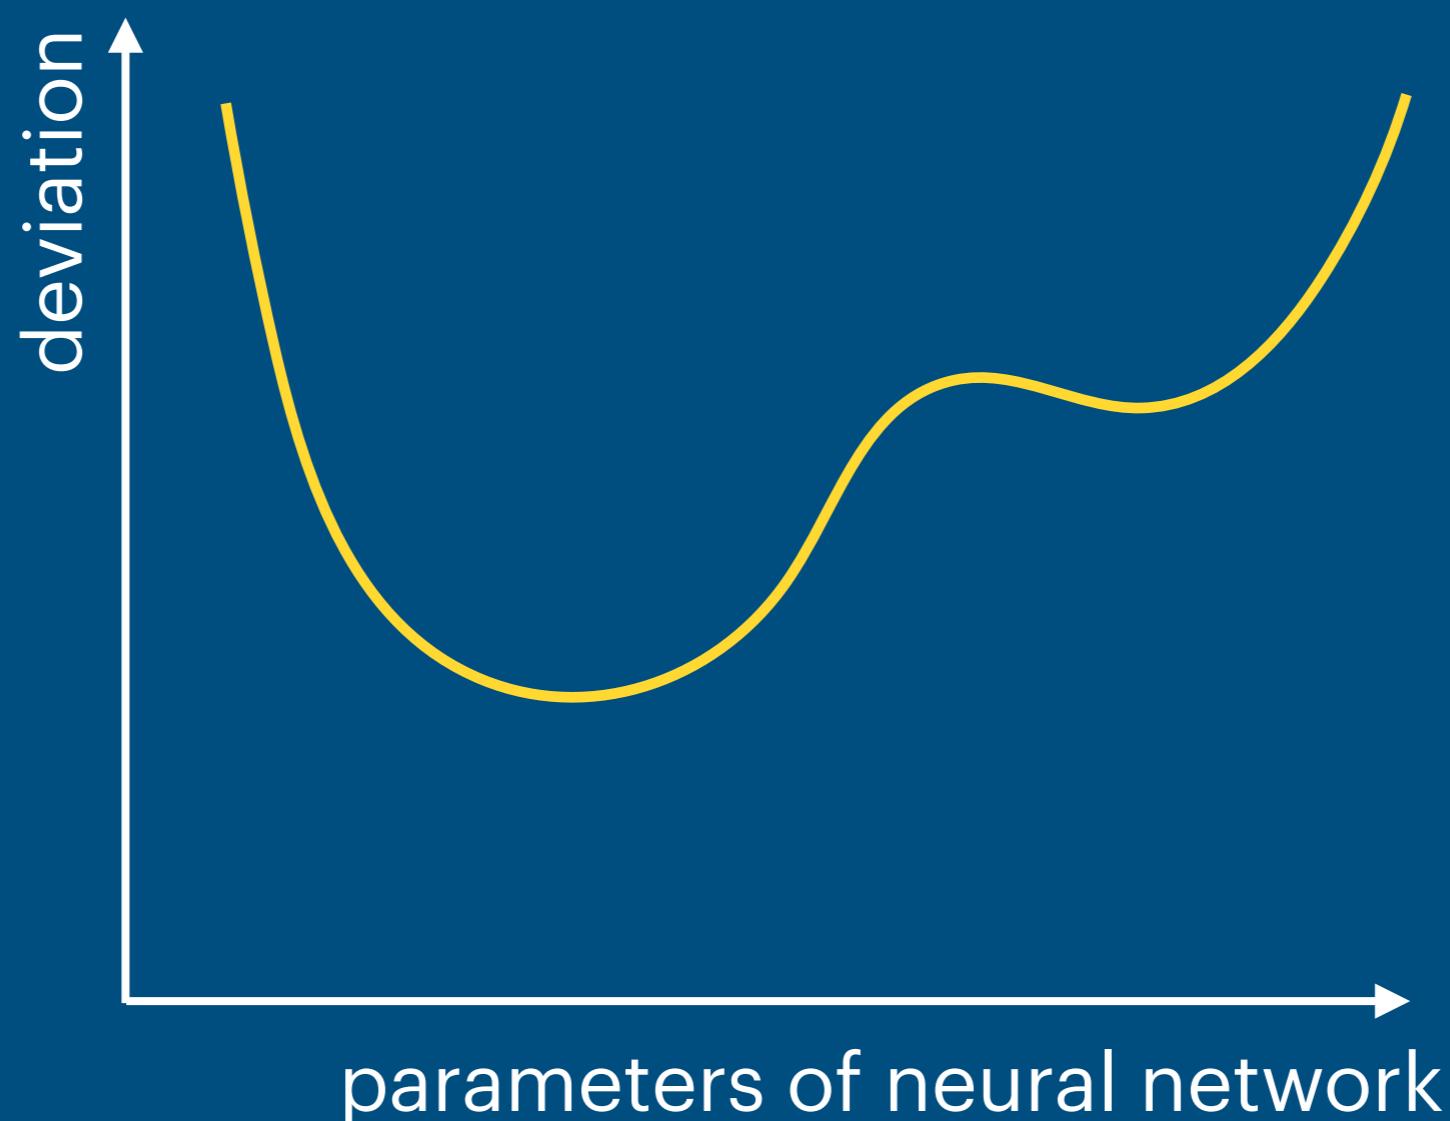
Everyone here will be able to apply neural networks to some research problem of their own (for example by starting from the code provided)

1

# Optimization

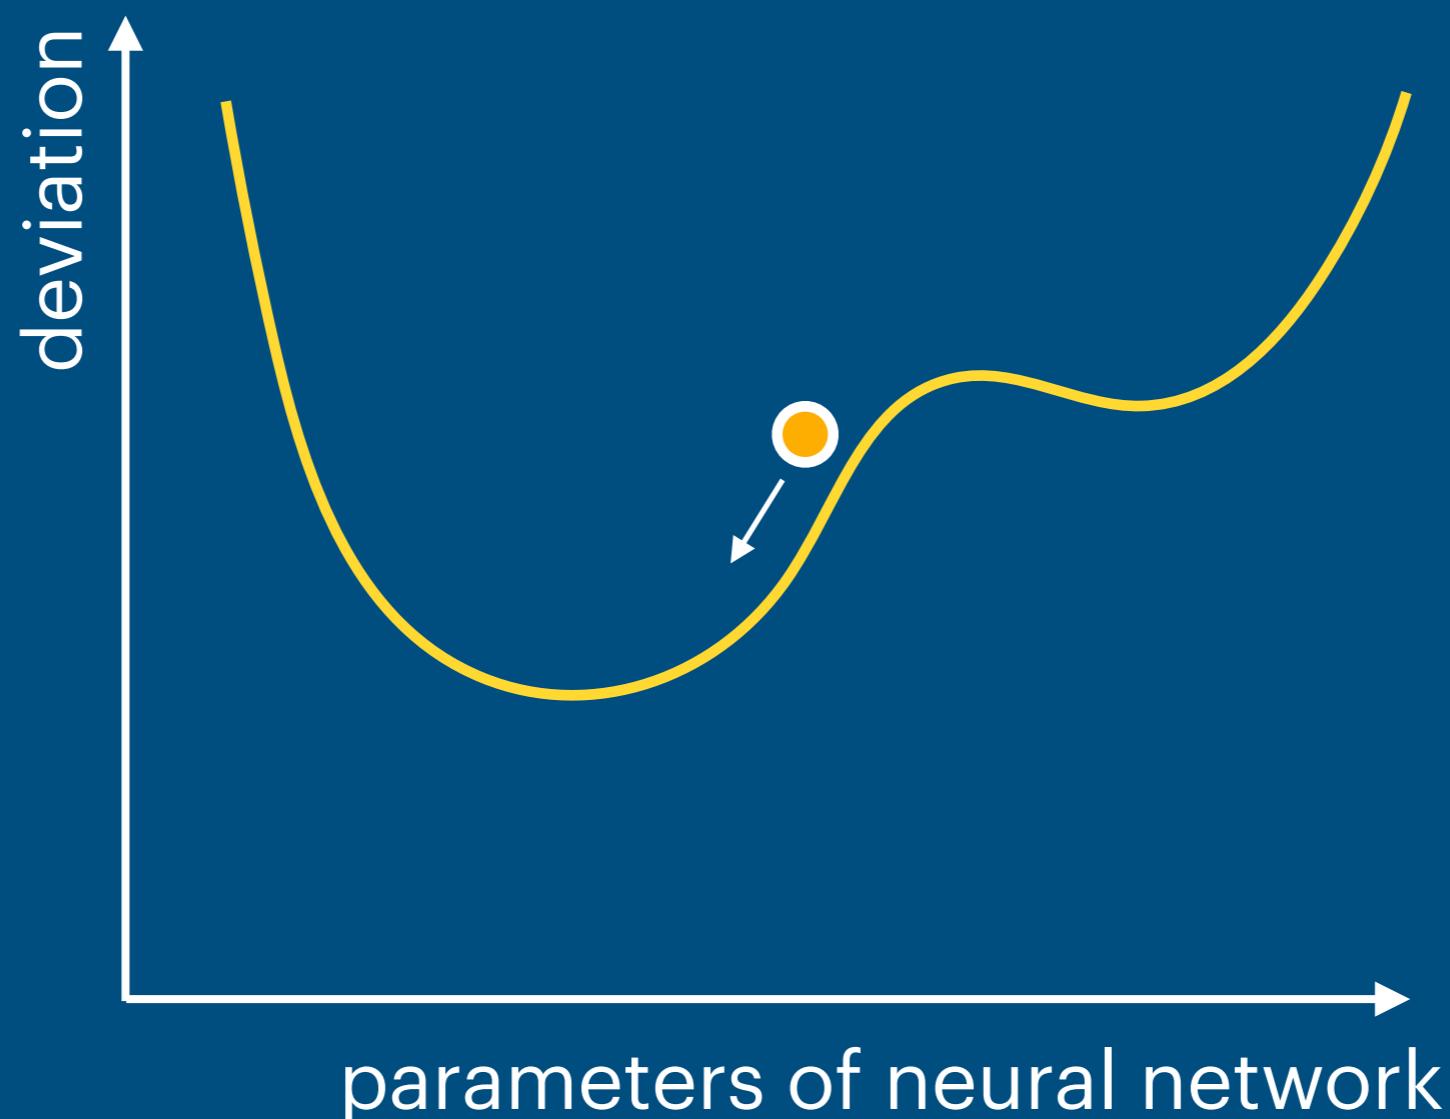
# Machine Learning = Optimization

Minimize the average deviation between the desired output of a neural network and its actual output !



# Optimization via Gradient Descent

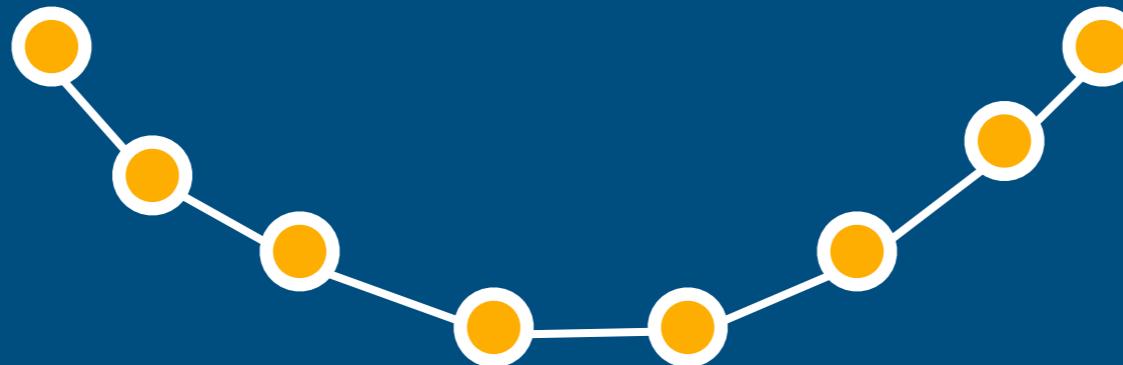
Walk down the hill  
(direction of steepest descent = negative gradient)



# Physics Example

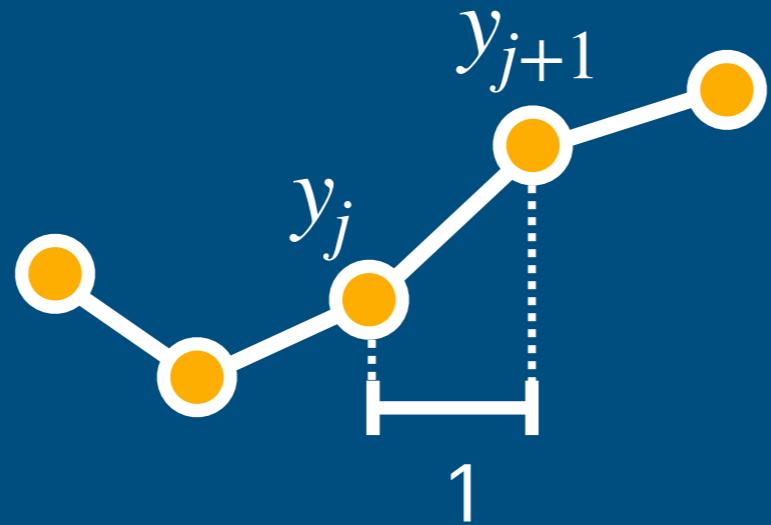
# Optimization via Gradient Descent

Physics example: minimize the energy of a mechanical structure to find its equilibrium shape!



Here: chain of masses, connected via springs, under gravity

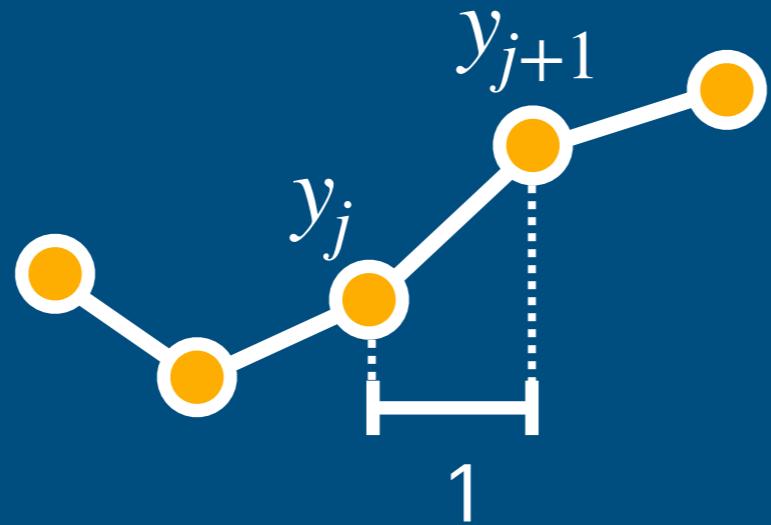
# A chain of masses with springs



Spring energy:

$$\sum_{j=0}^{N-1} \frac{K}{2} \left( \sqrt{(y_{j+1} - y_j)^2 + 1} - l \right)^2$$

# A chain of masses with springs



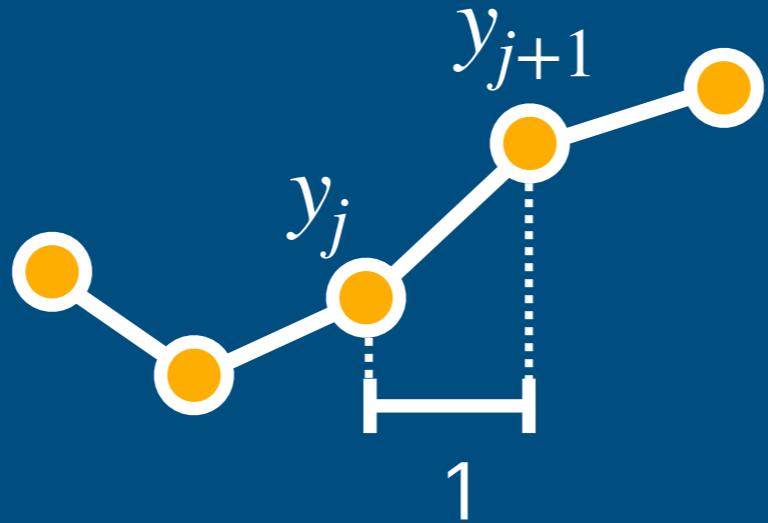
Spring energy:

$$\sum_{j=0}^{N-1} \frac{K}{2} \left( \sqrt{(y_{j+1} - y_j)^2 + 1} - l \right)^2$$

numpy:

```
(y[1:]-y[:-1])**2
```

## A chain of masses with springs



Spring energy:

$$\sum_{j=0}^{N-1} \frac{K}{2} \left( \sqrt{(y_{j+1} - y_j)^2 + 1} - l \right)^2$$

numpy:

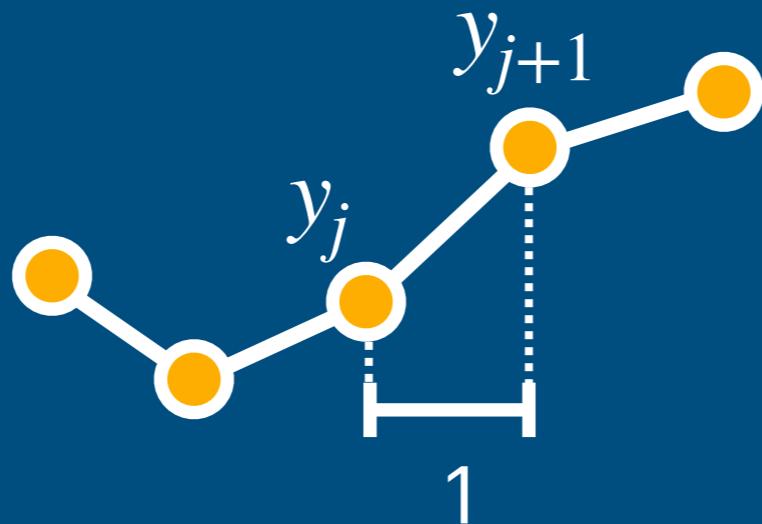
```
(y[1:]-y[:-1])**2
```

```
y = [8,7,9,5]
```

```
y[1:] = [7,9,5]
```

```
y[:-1] = [8,7,9]
```

## A chain of masses with springs

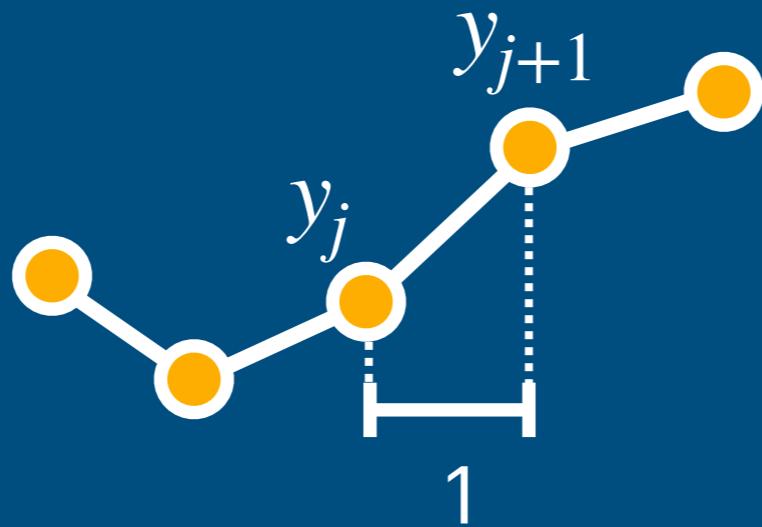


Spring energy:

$$\sum_{j=0}^{N-1} \frac{K}{2} \left( \sqrt{(y_{j+1} - y_j)^2 + 1} - l \right)^2$$

```
np.sum((np.sqrt((y[1:]-y[:-1])**2 + 1)-length)**2)
```

## A chain of masses with springs



```
def E(y, spring_constant, length):  
    ... calculate energy ...  
    return energy
```

...can also add gravity  $mgy_j$

# Minimize energy: Gradient descent



$$y_j^{\text{new}} = y_j^{\text{old}} - \eta \frac{\partial E}{\partial y_j}$$

↑  
step size  
(later: "learning rate")

# Minimize energy: Gradient descent



How to get the gradient?

Option 1: Calculate by hand and write function `grad_E`

Option 2: Use "automatic differentiation" library

# Interlude: jax

## Automatic Differentiation: **jax**

Take an arbitrary function and produce another function that calculates its gradient.

```
def f( x ):  
    return jnp.sum( x**2 )
```

$$f(x) = \sum_j x_j^2$$

```
def grad_f( x ):  
    return 2*x
```

$$\frac{\partial f}{\partial x_k} = 2x_k$$

# Automatic Differentiation: **jax**

Take an arbitrary function and produce another function that calculates its gradient.

```
def f( x ):  
    return jnp.sum( x**2 )
```



```
from jax import grad  
  
grad_f = grad( f )
```



...is equivalent to:

```
def grad_f( x ):  
    return 2*x
```

$$f(x) = \sum_j x_j^2$$

$$\frac{\partial f}{\partial x_k} = 2x_k$$

## Two steps to use jax

(1) Instead of numpy use jax.numpy for everything:

```
import jax.numpy as jnp
```

```
y = jnp.sin( x )
```

Avoid for loops when possible, for efficiency.

Use arrays for everything.

No conditionals: replace

```
if x<0:      by    y = (x<0) * (-x) + (x>=0) * (x**2)
    y=-x
else:
    y=x**2
```

## Two steps to use jax

### (2) Functional programming:

All program functions must be like mathematical functions

No side effects:

The function does not change its arguments  
and does not change or rely on any global variables etc.  
and does not print anything

This makes it much easier to understand the  
flow of information!

jax is simple but state of the art

jax is simple to use  
(numpy + grad + handful of other simple but powerful commands)

jax has very good performance

jax is used in state-of-the-art projects, like DeepMind's AlphaFold

main alternatives: tensorflow, pytorch

# Back to physics

# Minimize energy: Gradient descent



$$y_j^{\text{new}} = y_j^{\text{old}} - \eta \frac{\partial E}{\partial y_j}$$

↑  
step size  
(later: "learning rate")

## Minimize energy: Gradient descent



Produce new function that calculates the gradient of E:

```
grad_E = grad(E, argnum=0)
```

(with respect to its first argument  $y$ , indicated by  $\text{argnum}=0$ )

Now do gradient descent:

```
y = y - eta * grad_E(y, spring, length, ...)
```

# Let us try it out!

[FlorianMarquardt.github.io/  
MachineLearningThreeEasyLessons](https://FlorianMarquardt.github.io/MachineLearningThreeEasyLessons)

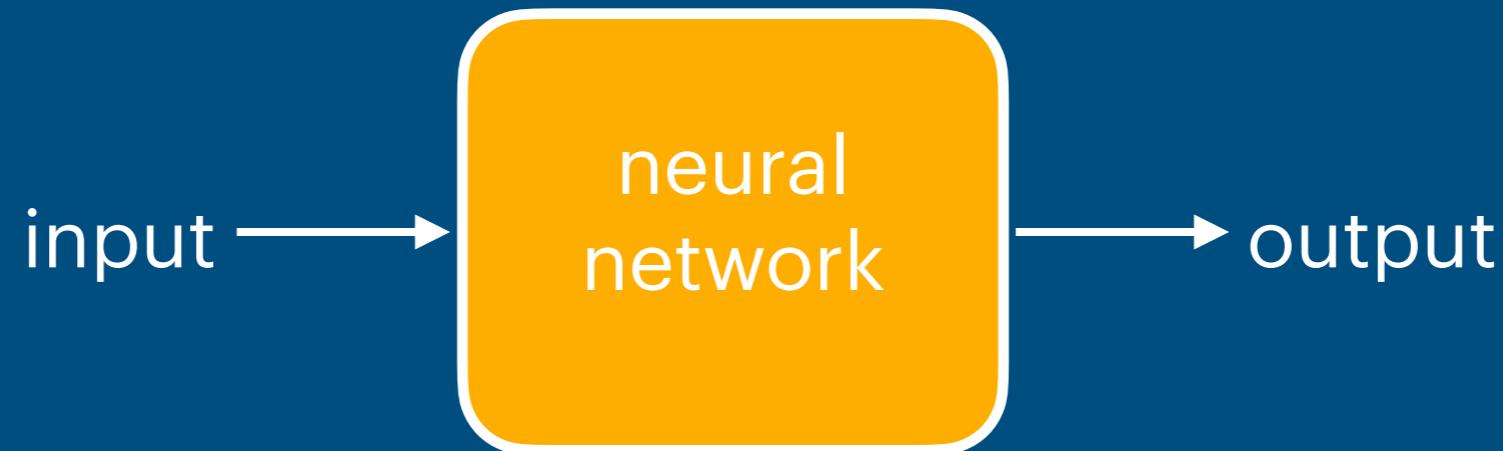
(Lesson 1)

2

# Neural Networks

# Machine Learning = Optimization

Minimize the average deviation between the desired output of a neural network and its actual output !

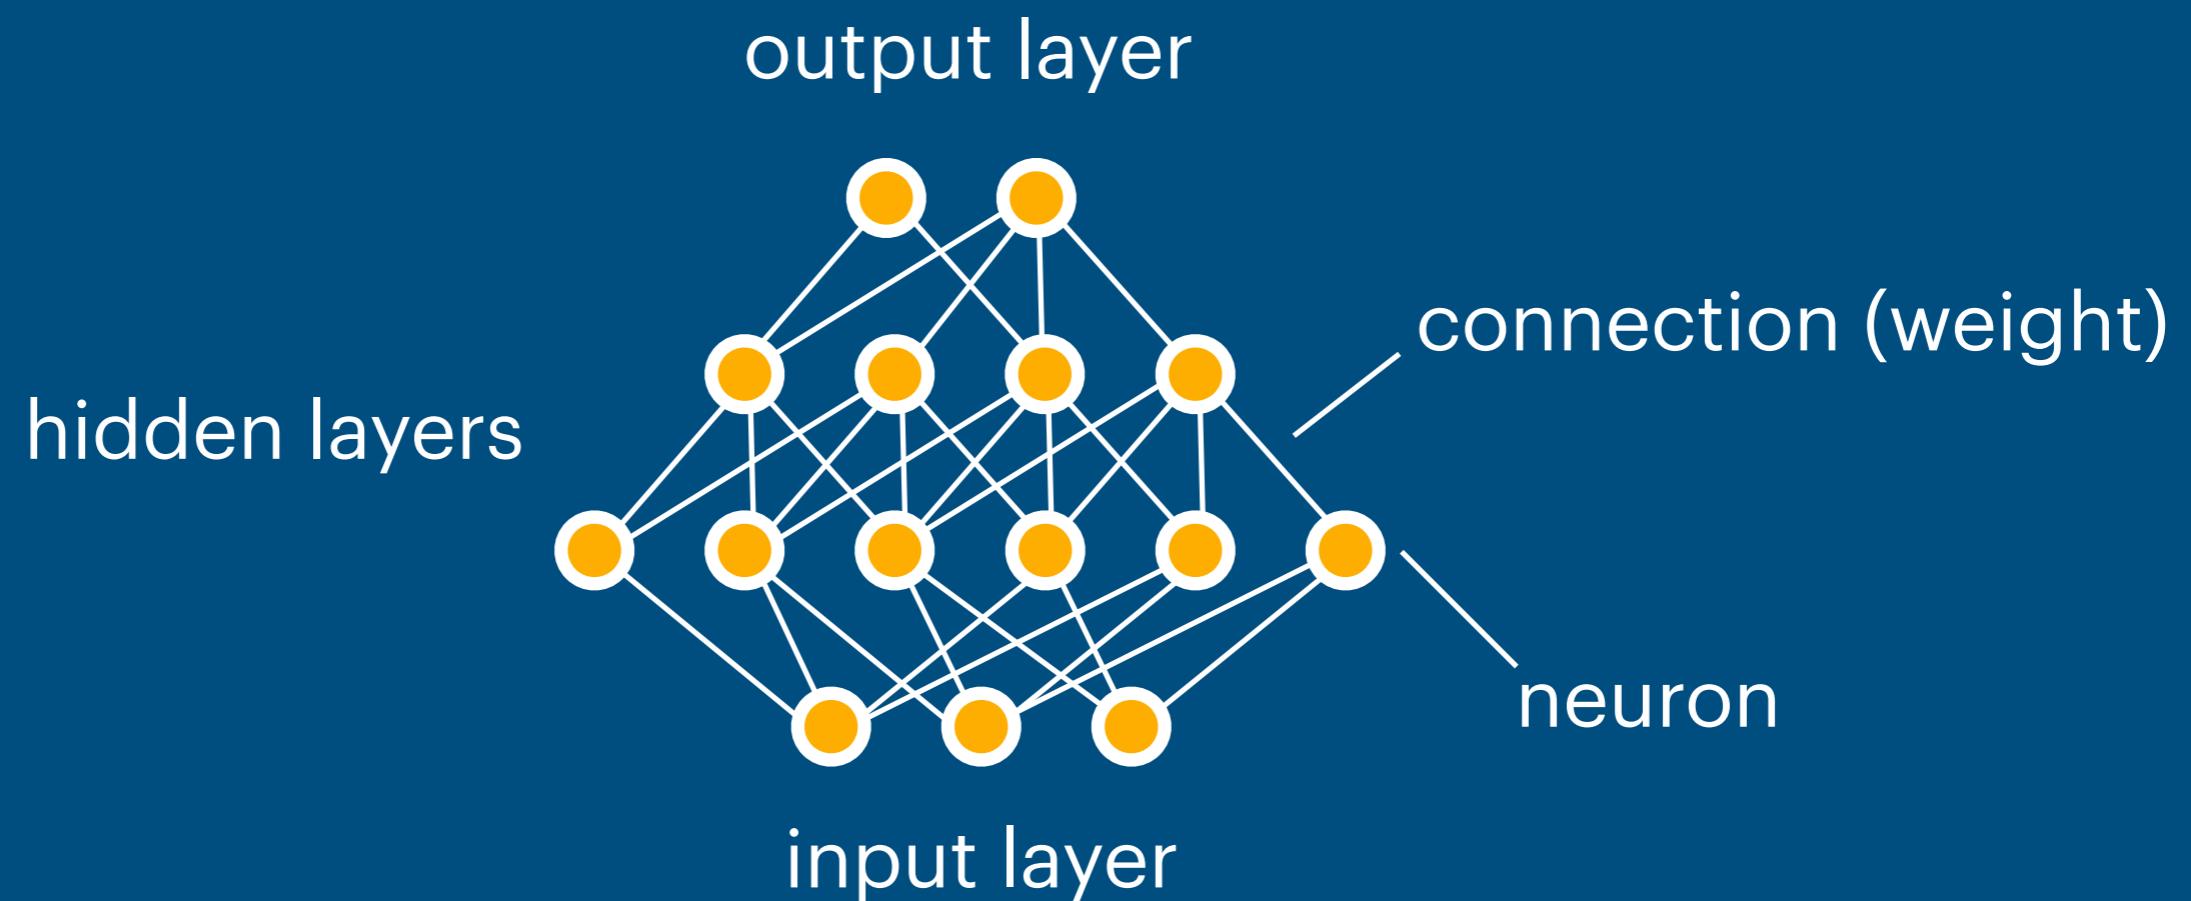


"loss" (or "cost") = average of  $(\text{target output} - \text{actual output})^2$   
averaged over all possible inputs

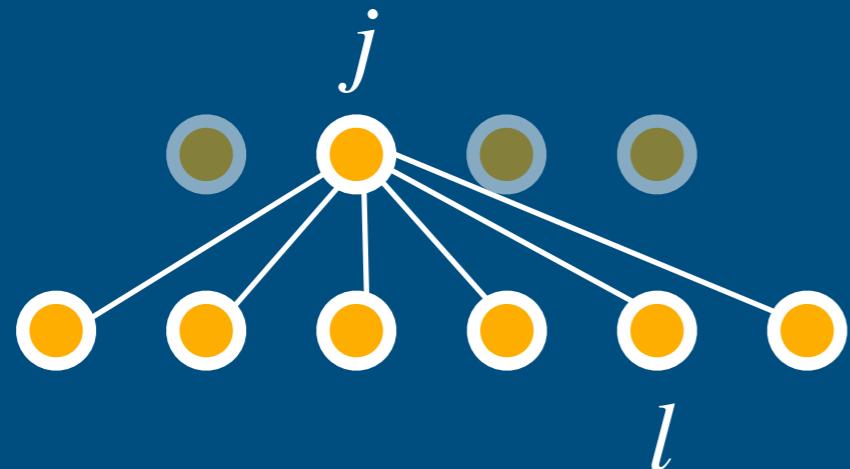
# Neural network structure

Artificial neural network = high-dimensional function depending on many parameters

Artificial neural network = cartoon of biological neural network



# Neural network evaluation



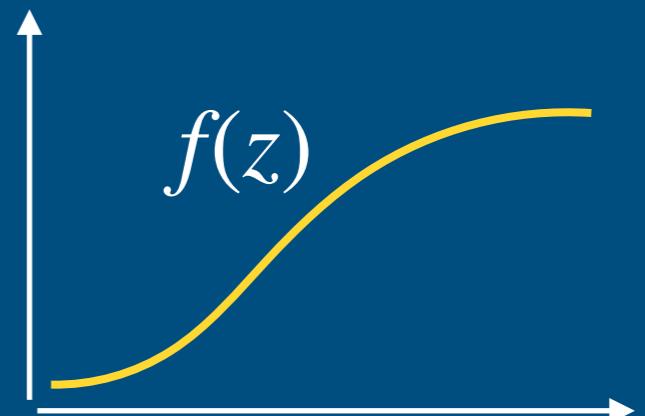
$x$  = vector of neuron activations in lower layer

$$(1) \text{ calculate } z_j = \sum_l w_{jl} x_l + b_j$$

weights      biases

$$(2) \text{ and then } x_j^{\text{new}} = f(z_j)$$

nonlinear activation function  $f$ :



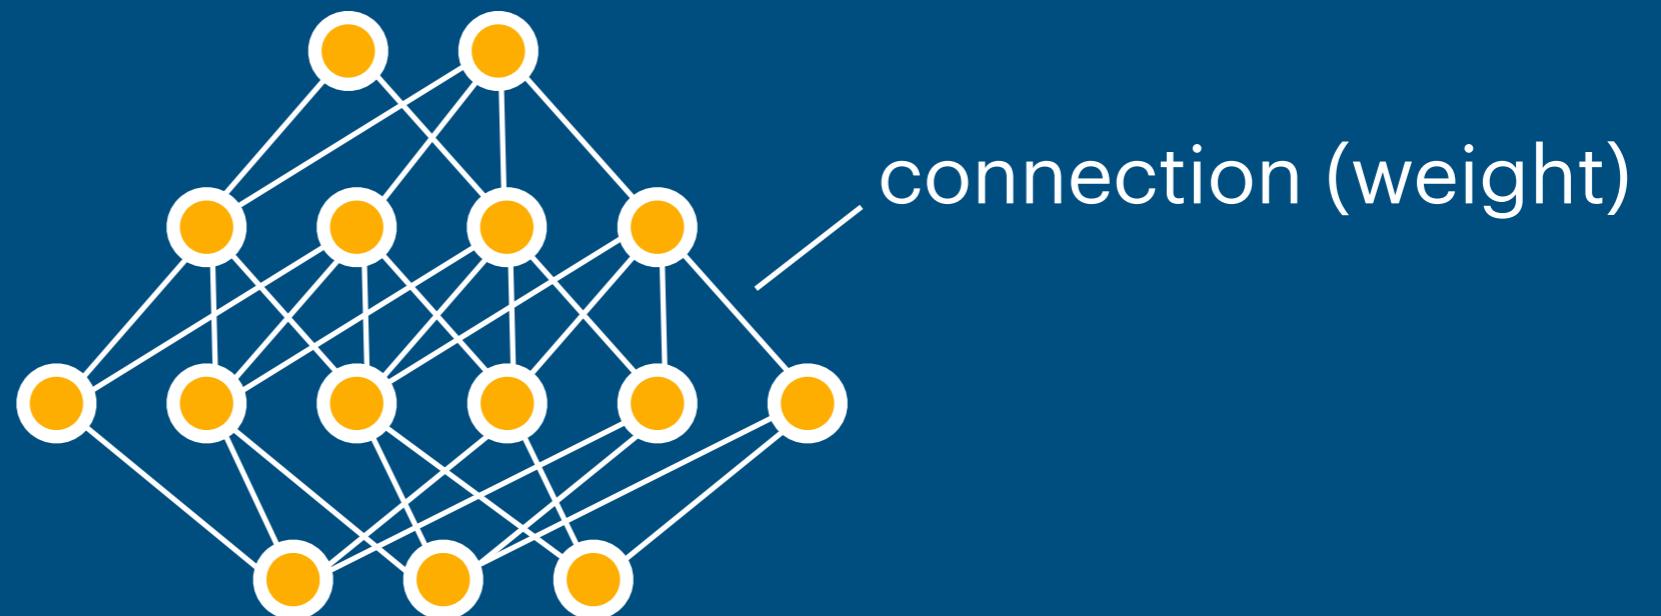
$x^{\text{new}}$  = vector of neuron activations in upper layer

# Neural network evaluation

layer by layer:

$$\text{calculate } z_j = \sum_l w_{jl}x_l + b_j \quad \text{and then } x_j^{\text{new}} = f(z_j)$$

w: weight matrix, b: bias vector (different for each layer)



# Neural network evaluation

layer by layer:

$$\text{calculate } z_j = \sum_l w_{jl}x_l + b_j \quad \text{and then } x_j^{\text{new}} = f(z_j)$$

w: weight matrix, b: bias vector (different for each layer)

```
x = jnp.matmul( w, x) + b
```

```
x = f( x )
```

# Neural network evaluation

layer by layer:

$$\text{calculate } z_j = \sum_l w_{jl}x_l + b_j \quad \text{and then } x_j^{\text{new}} = f(z_j)$$

w: weight matrix, b: bias vector (different for each layer)

Simplest neural network:

```
def NN( x, weights, biases ) :
    for w,b in zip(weights,biases) :
        x = jnp.matmul( w, x ) + b
        x = f( x )
    return x
```

weights: list of matrices

biases: list of vectors

# Neural network evaluation

Slightly nicer:

```
def NN( x, params ):  
    for w,b in zip(params['weights'],params['biases']):  
        x = jnp.matmul( w, x ) + b  
    x = f( x )  
return x
```

params: a dictionary containing the weights and biases

## Loss function

simplest choice: "mean square error"

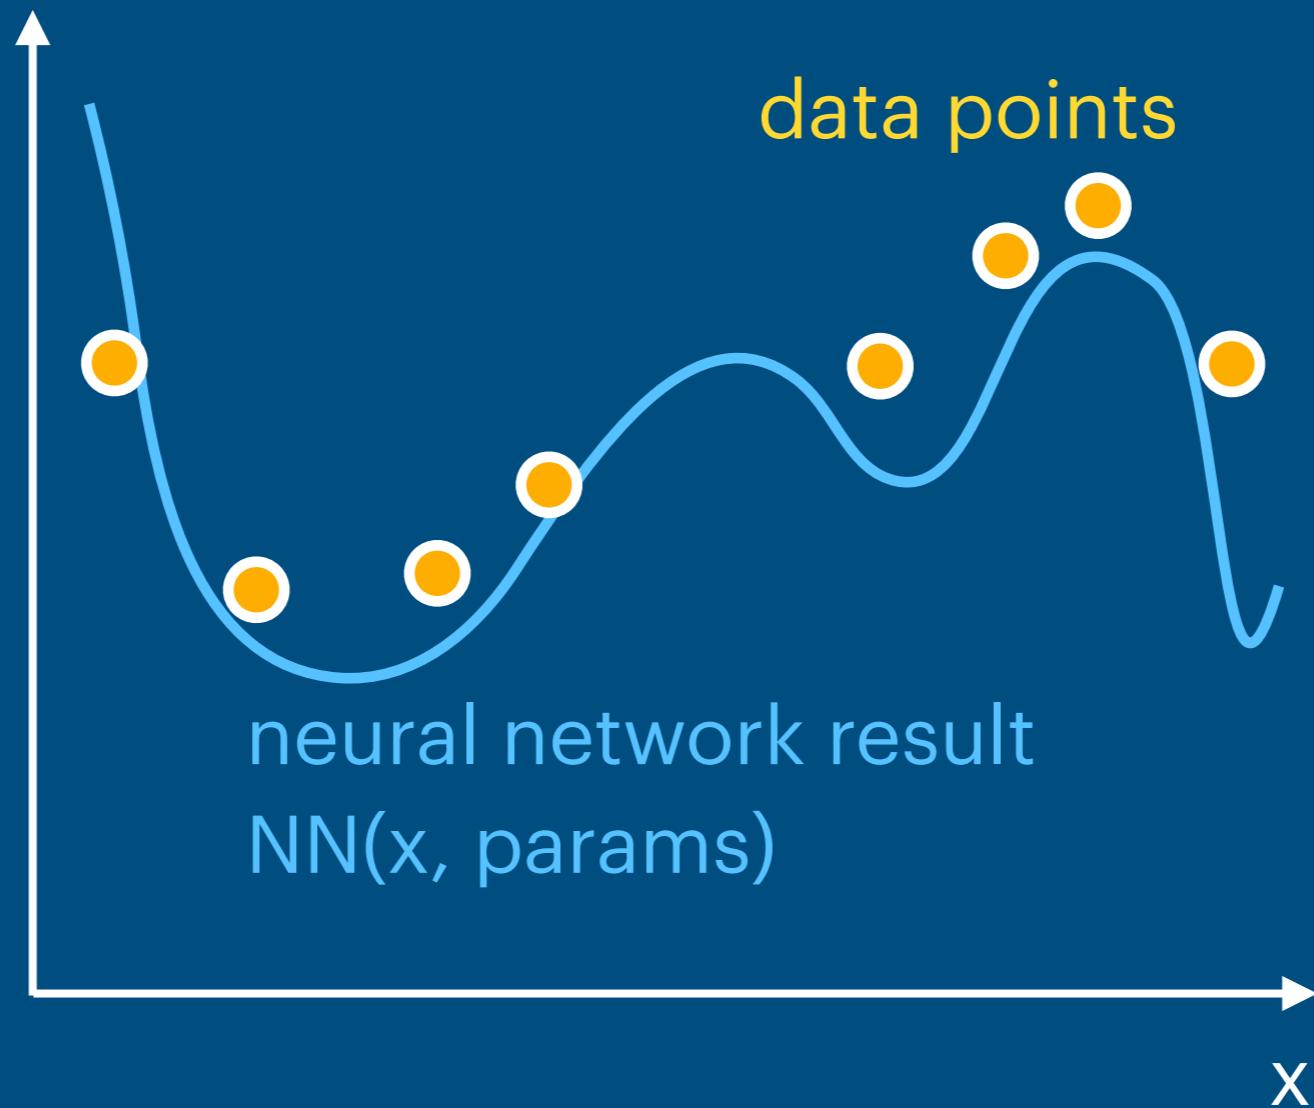
```
def mse_loss( x, y_target, params ):  
    return jnp.sum( (NN( x, params ) - y_target )**2 )
```

Important: this calls the network, and has an explicit dependence on the parameters **params**

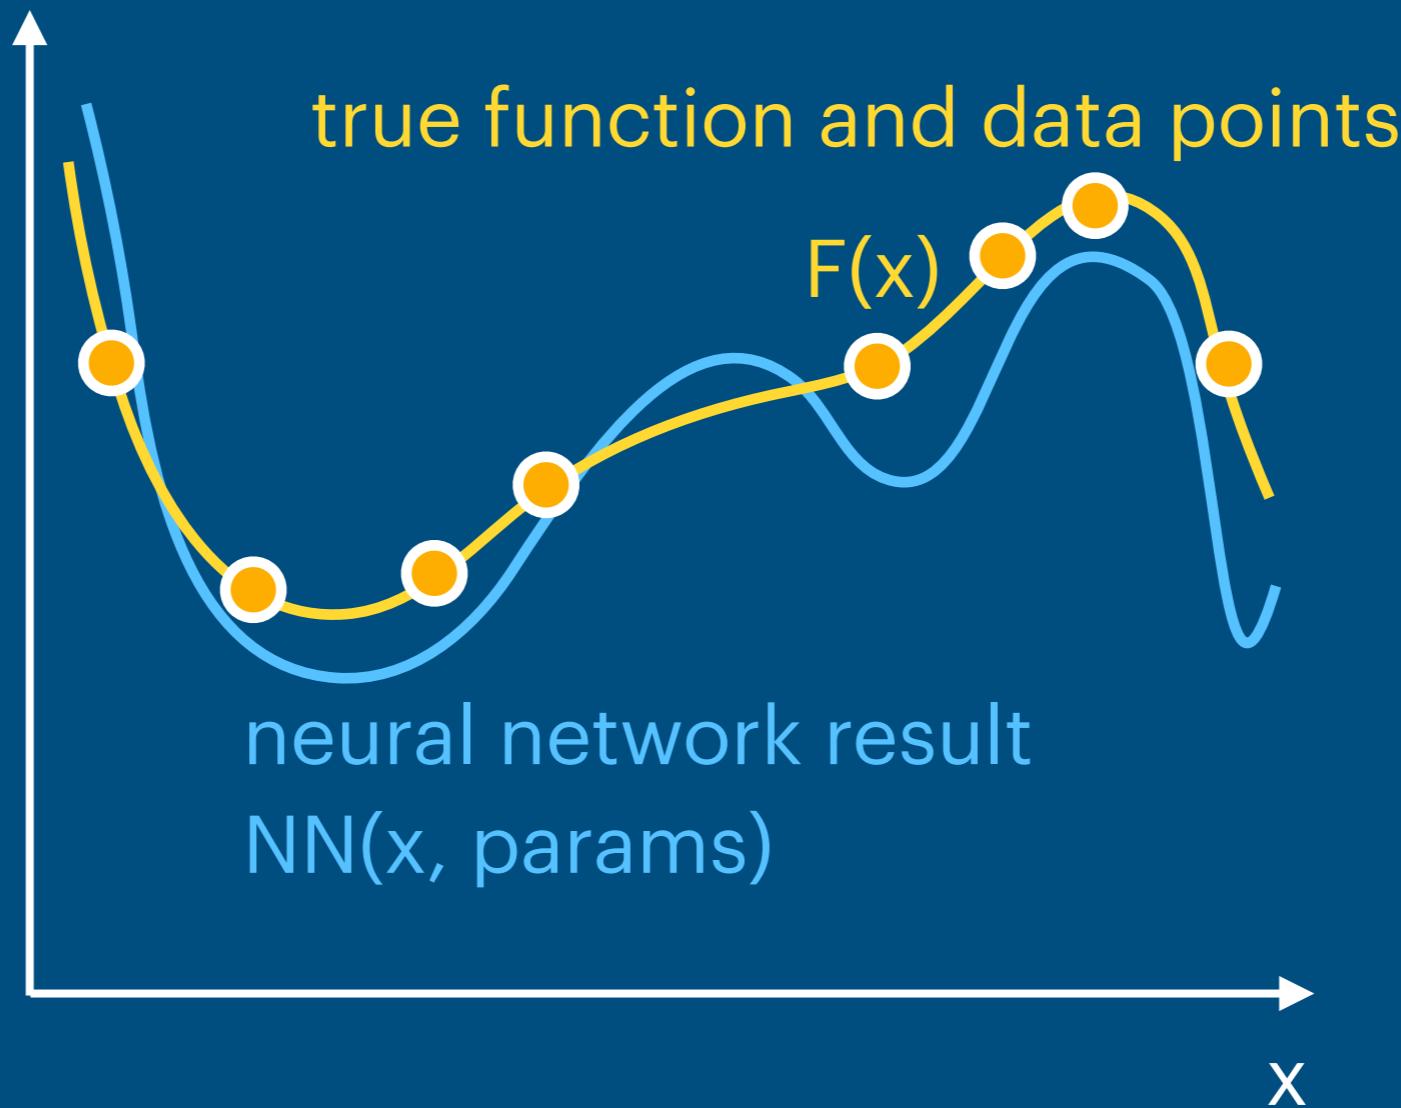
Thus, we can take the gradient with respect to the parameters!

```
mse_loss_grad = grad( mse_loss, argnums=2 )
```

# An example: fitting data / a function via a neural network



## An example: fitting data / a function via a neural network



In our example: sample random  $x$  positions, evaluate true function  $F(x)$  to get the data points, and try to minimize mean square error for the network.

## Two more things

- jax random numbers (see code)
- batch processing (see code)  
processing `x` of shape `[batchsize, input_dim]` in parallel

# Let us try it out!

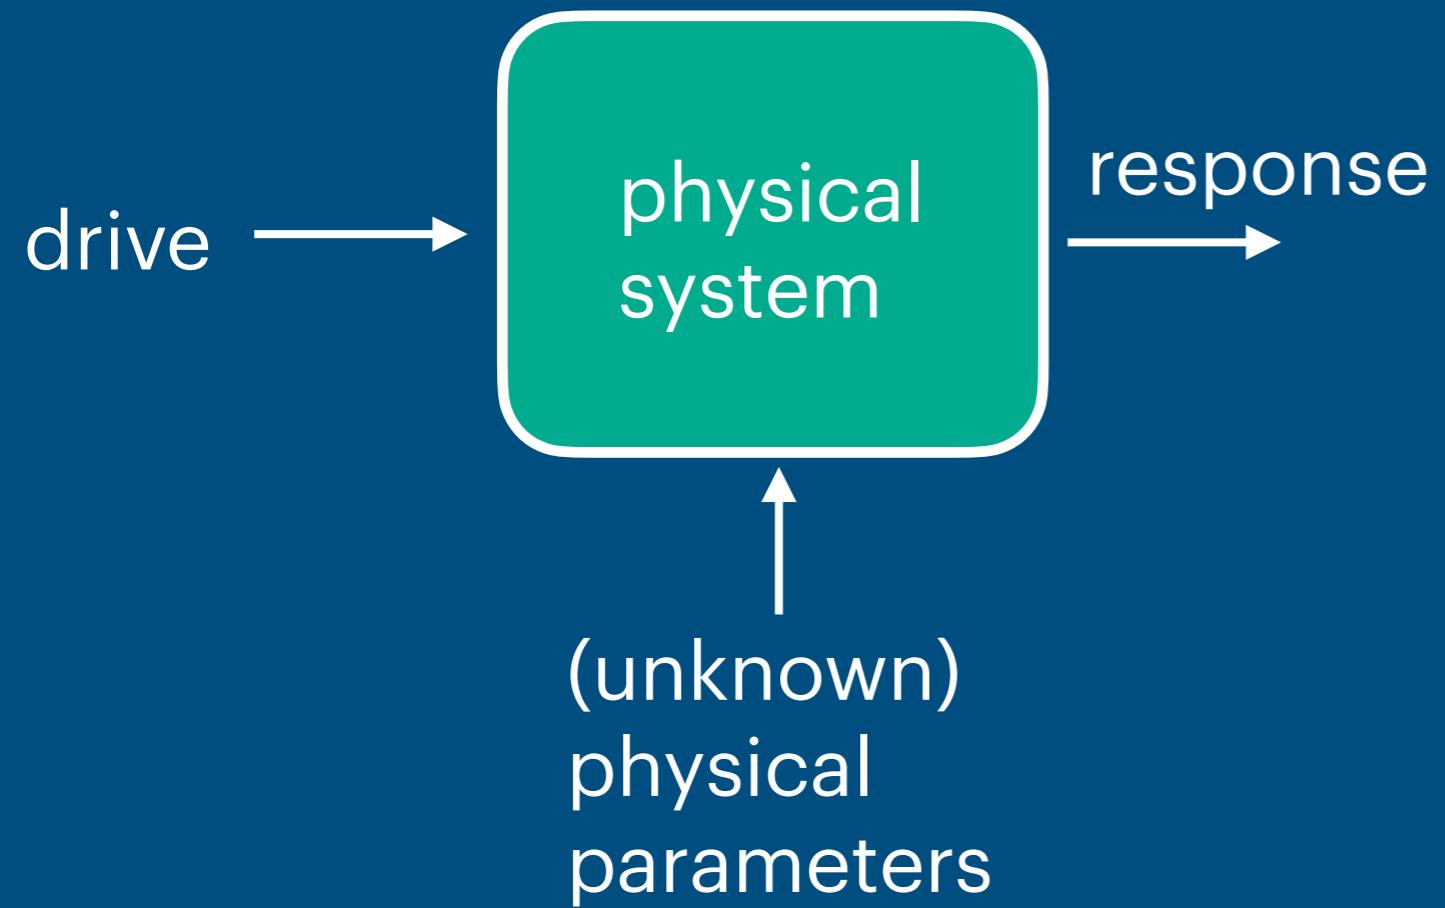
[FlorianMarquardt.github.io/  
MachineLearningThreeEasyLessons](https://FlorianMarquardt.github.io/MachineLearningThreeEasyLessons)

(Lesson 2)

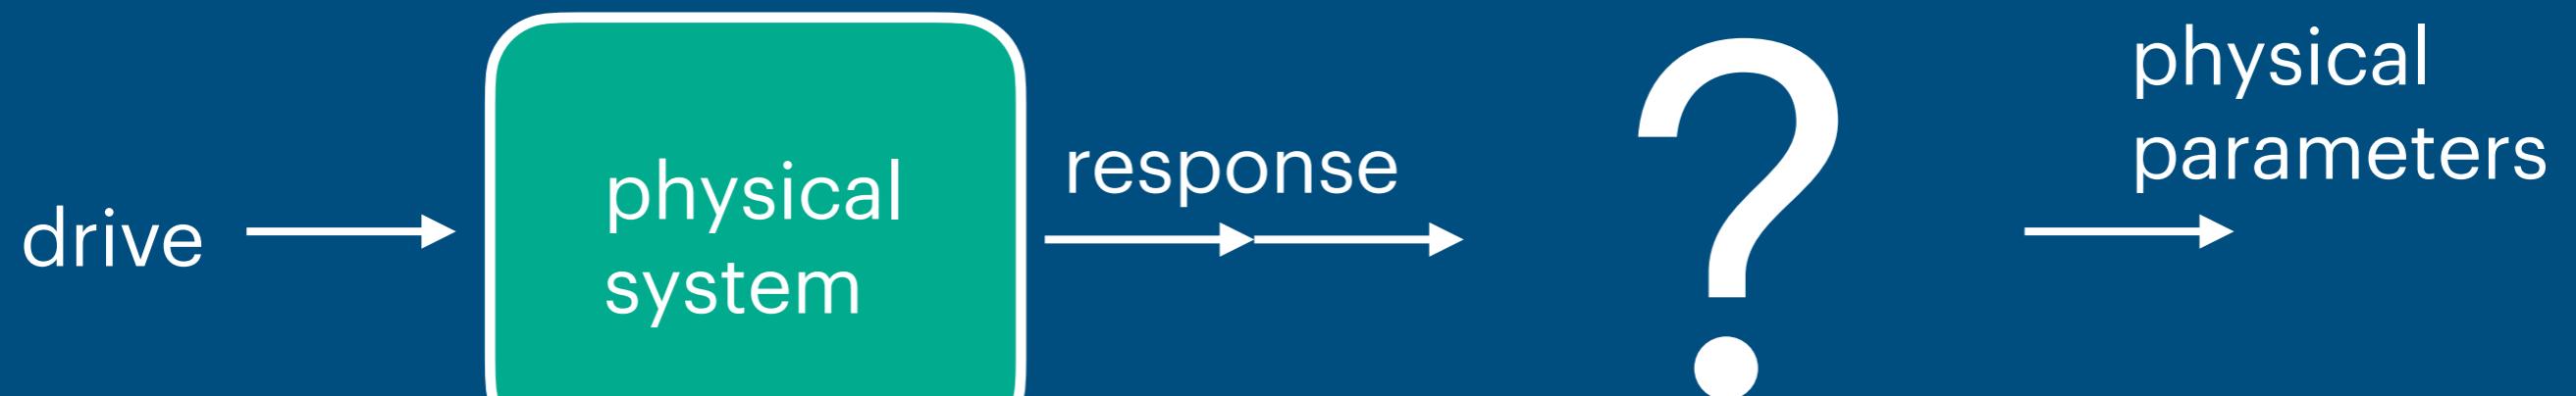
3

# Full example: Characterizing a nonlinear oscillator

# Physics example: Characterizing a physical system



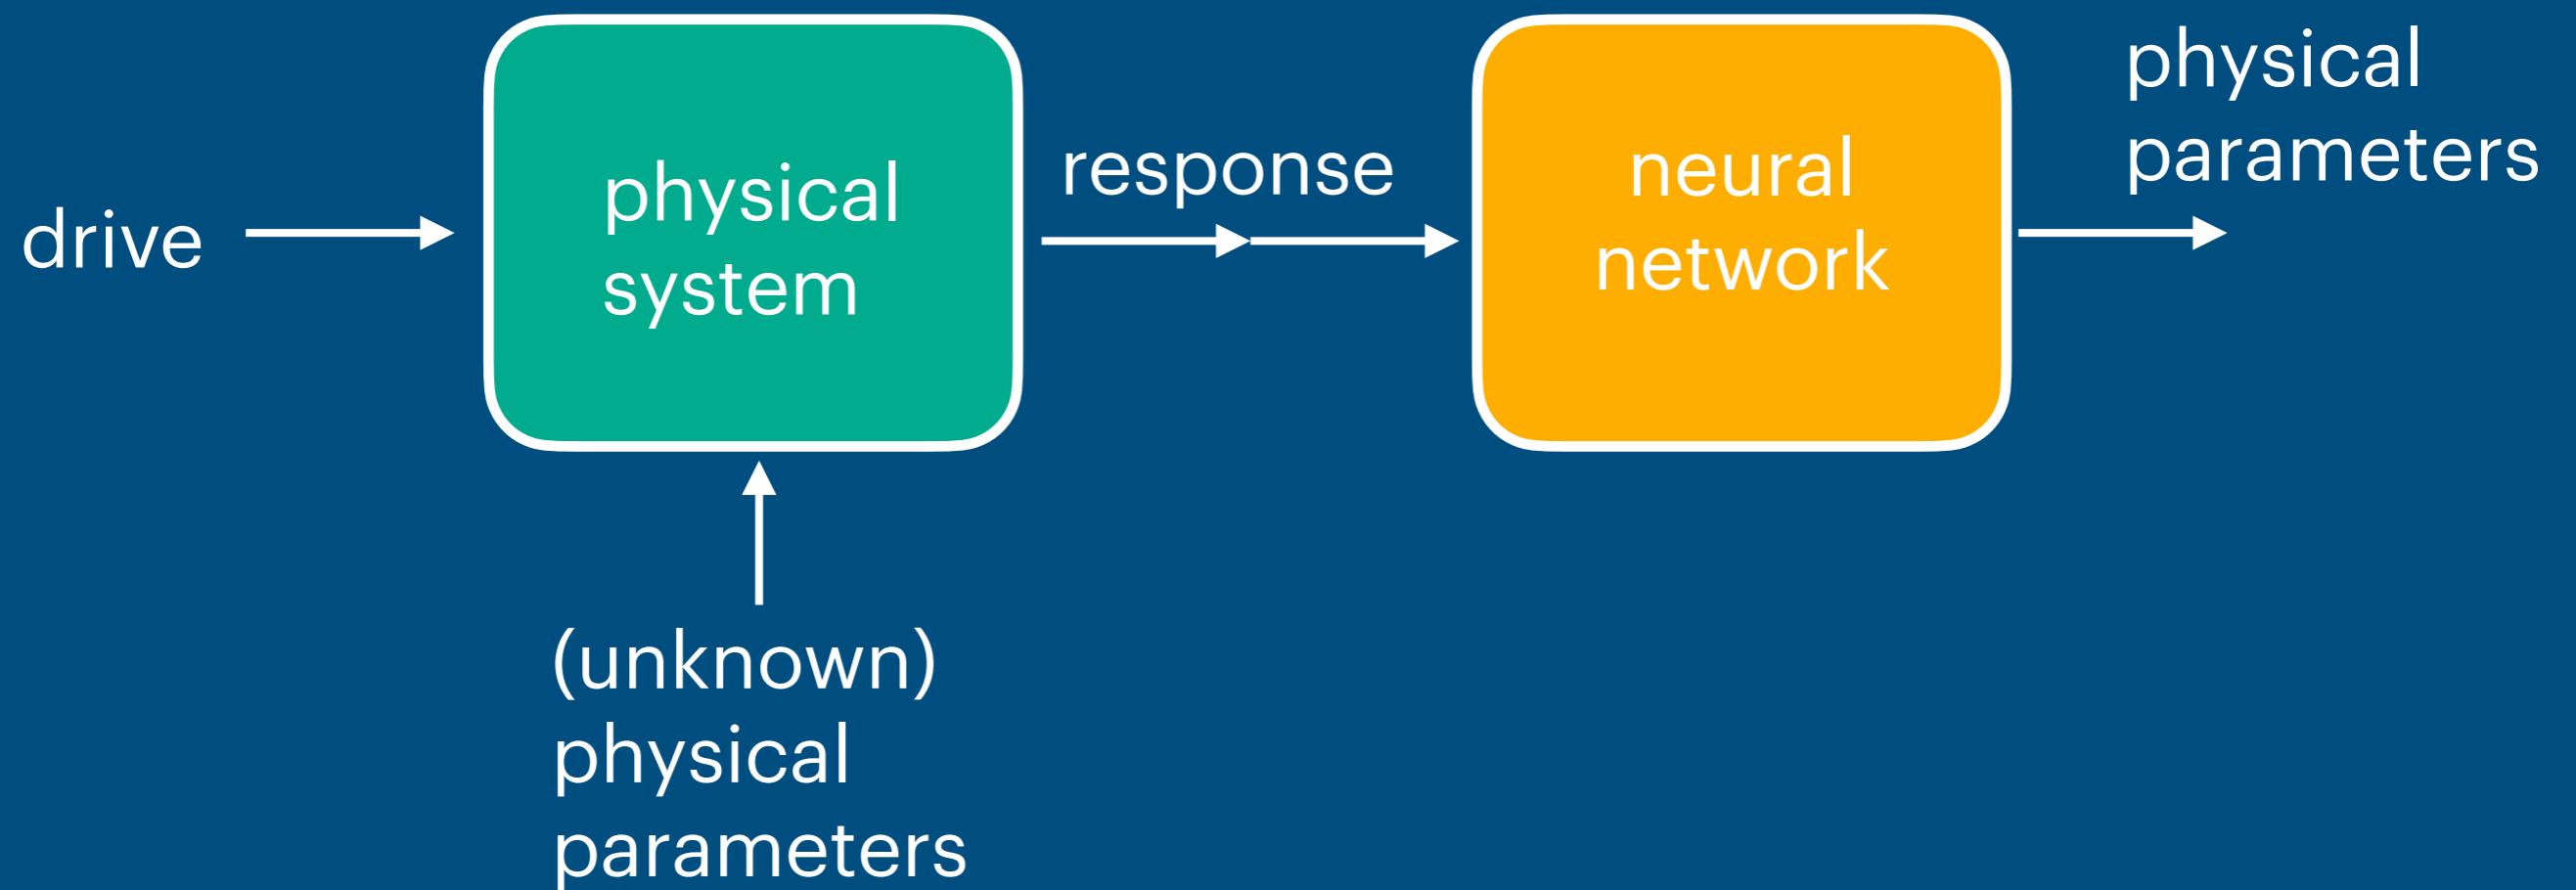
# Physics example: Characterizing a physical system



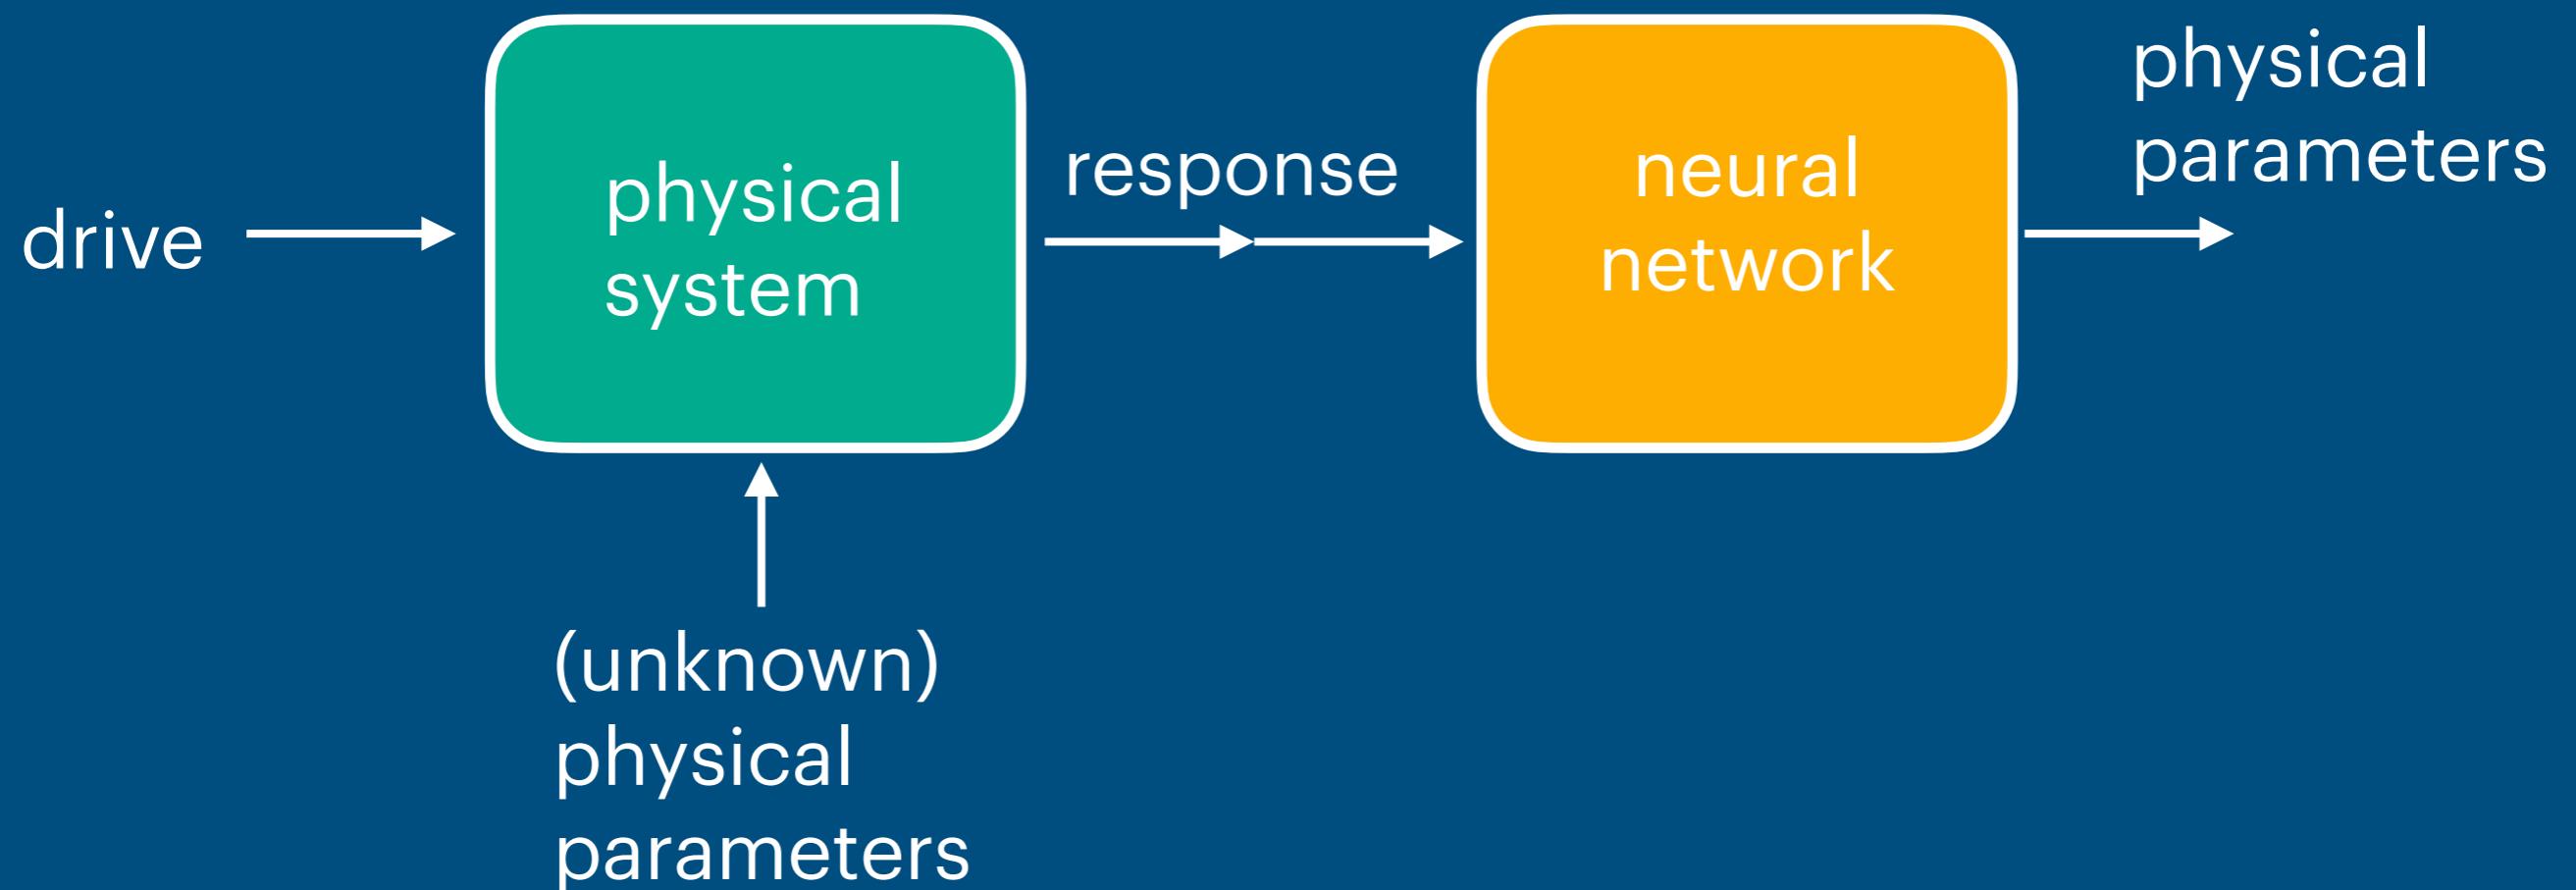
Difficult **inverse** problem. Usual options:

- (1) if response is known analytically:  
nonlinear curve fitting with gradient descent in physical parameters  
(but: slow, may get stuck)
- (2) if it is not known analytically, but simulations available: gradient-free optimization of phys. parameters  
(even slower)
- (3) if there are not even simulations:  
practically impossible

# Physics example: Characterizing a physical system

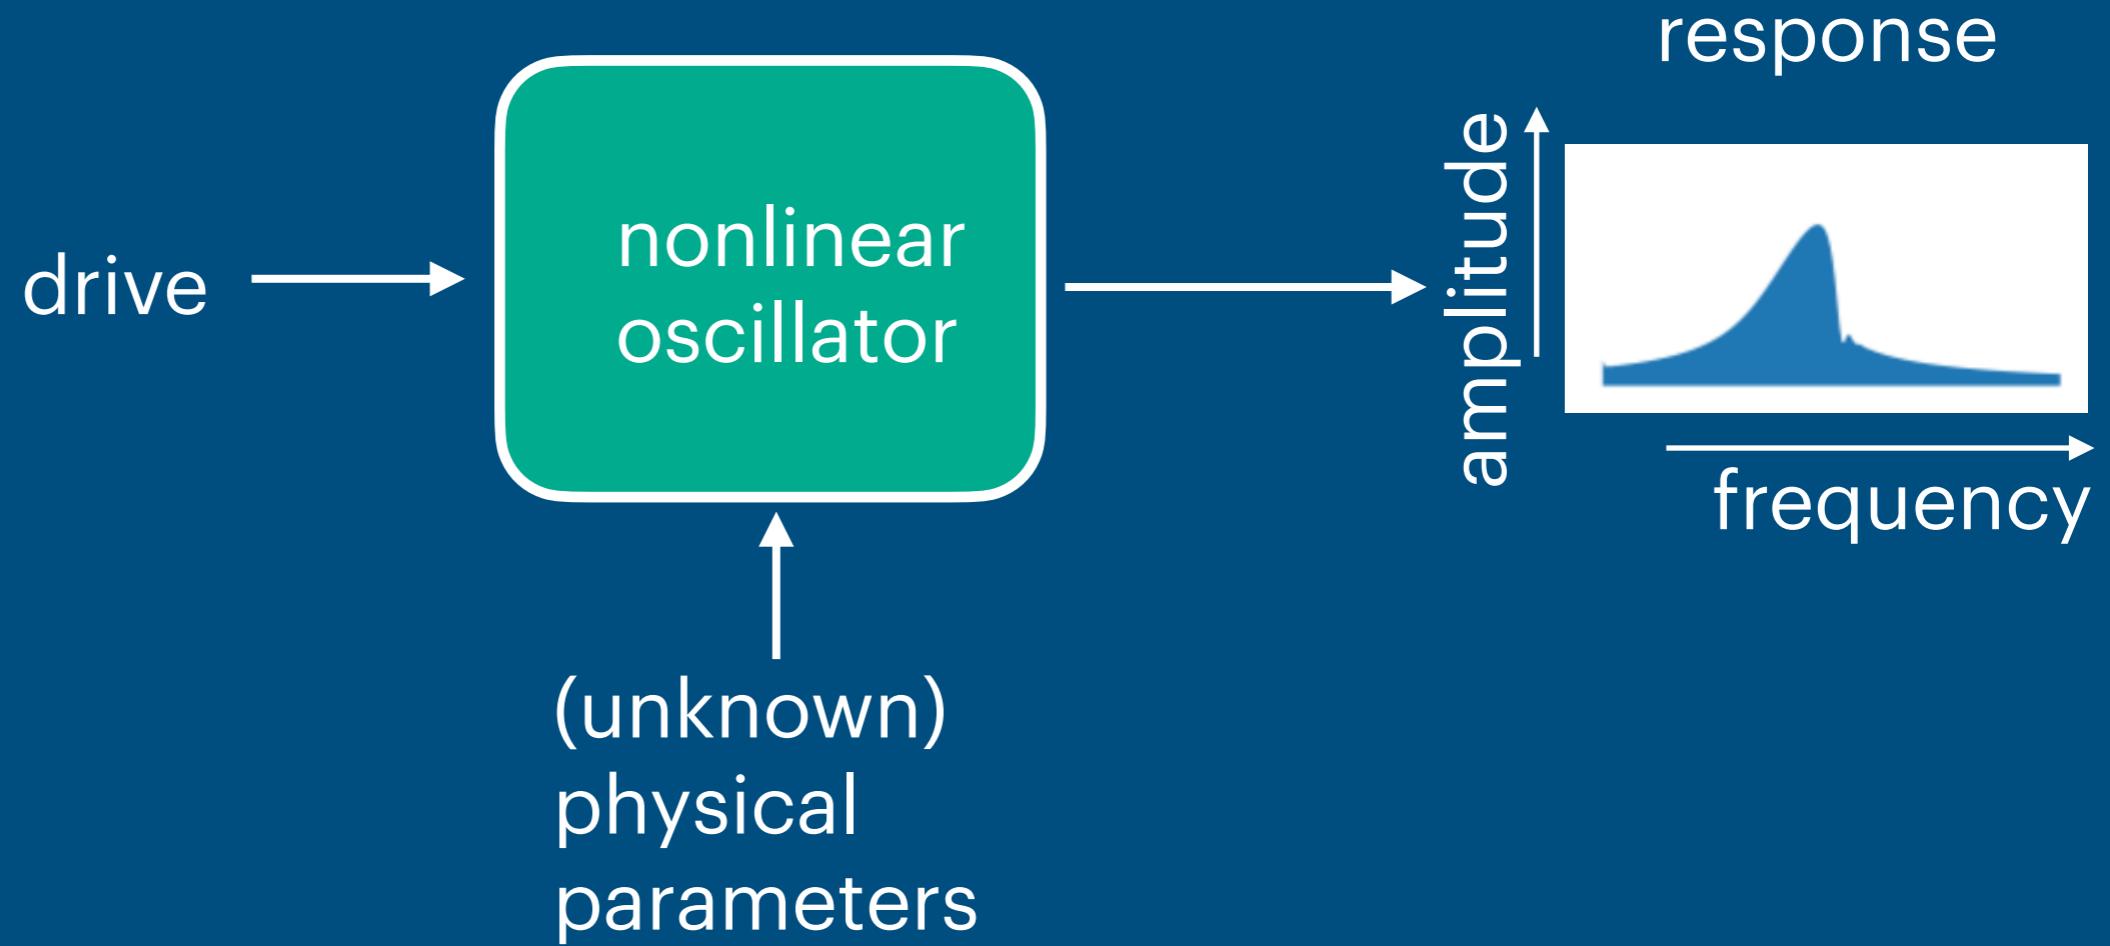


# Physics example: Characterizing a physical system



train on many examples where true  
parameters are known!

## Physics example: Characterizing a physical system



# Duffing oscillator:

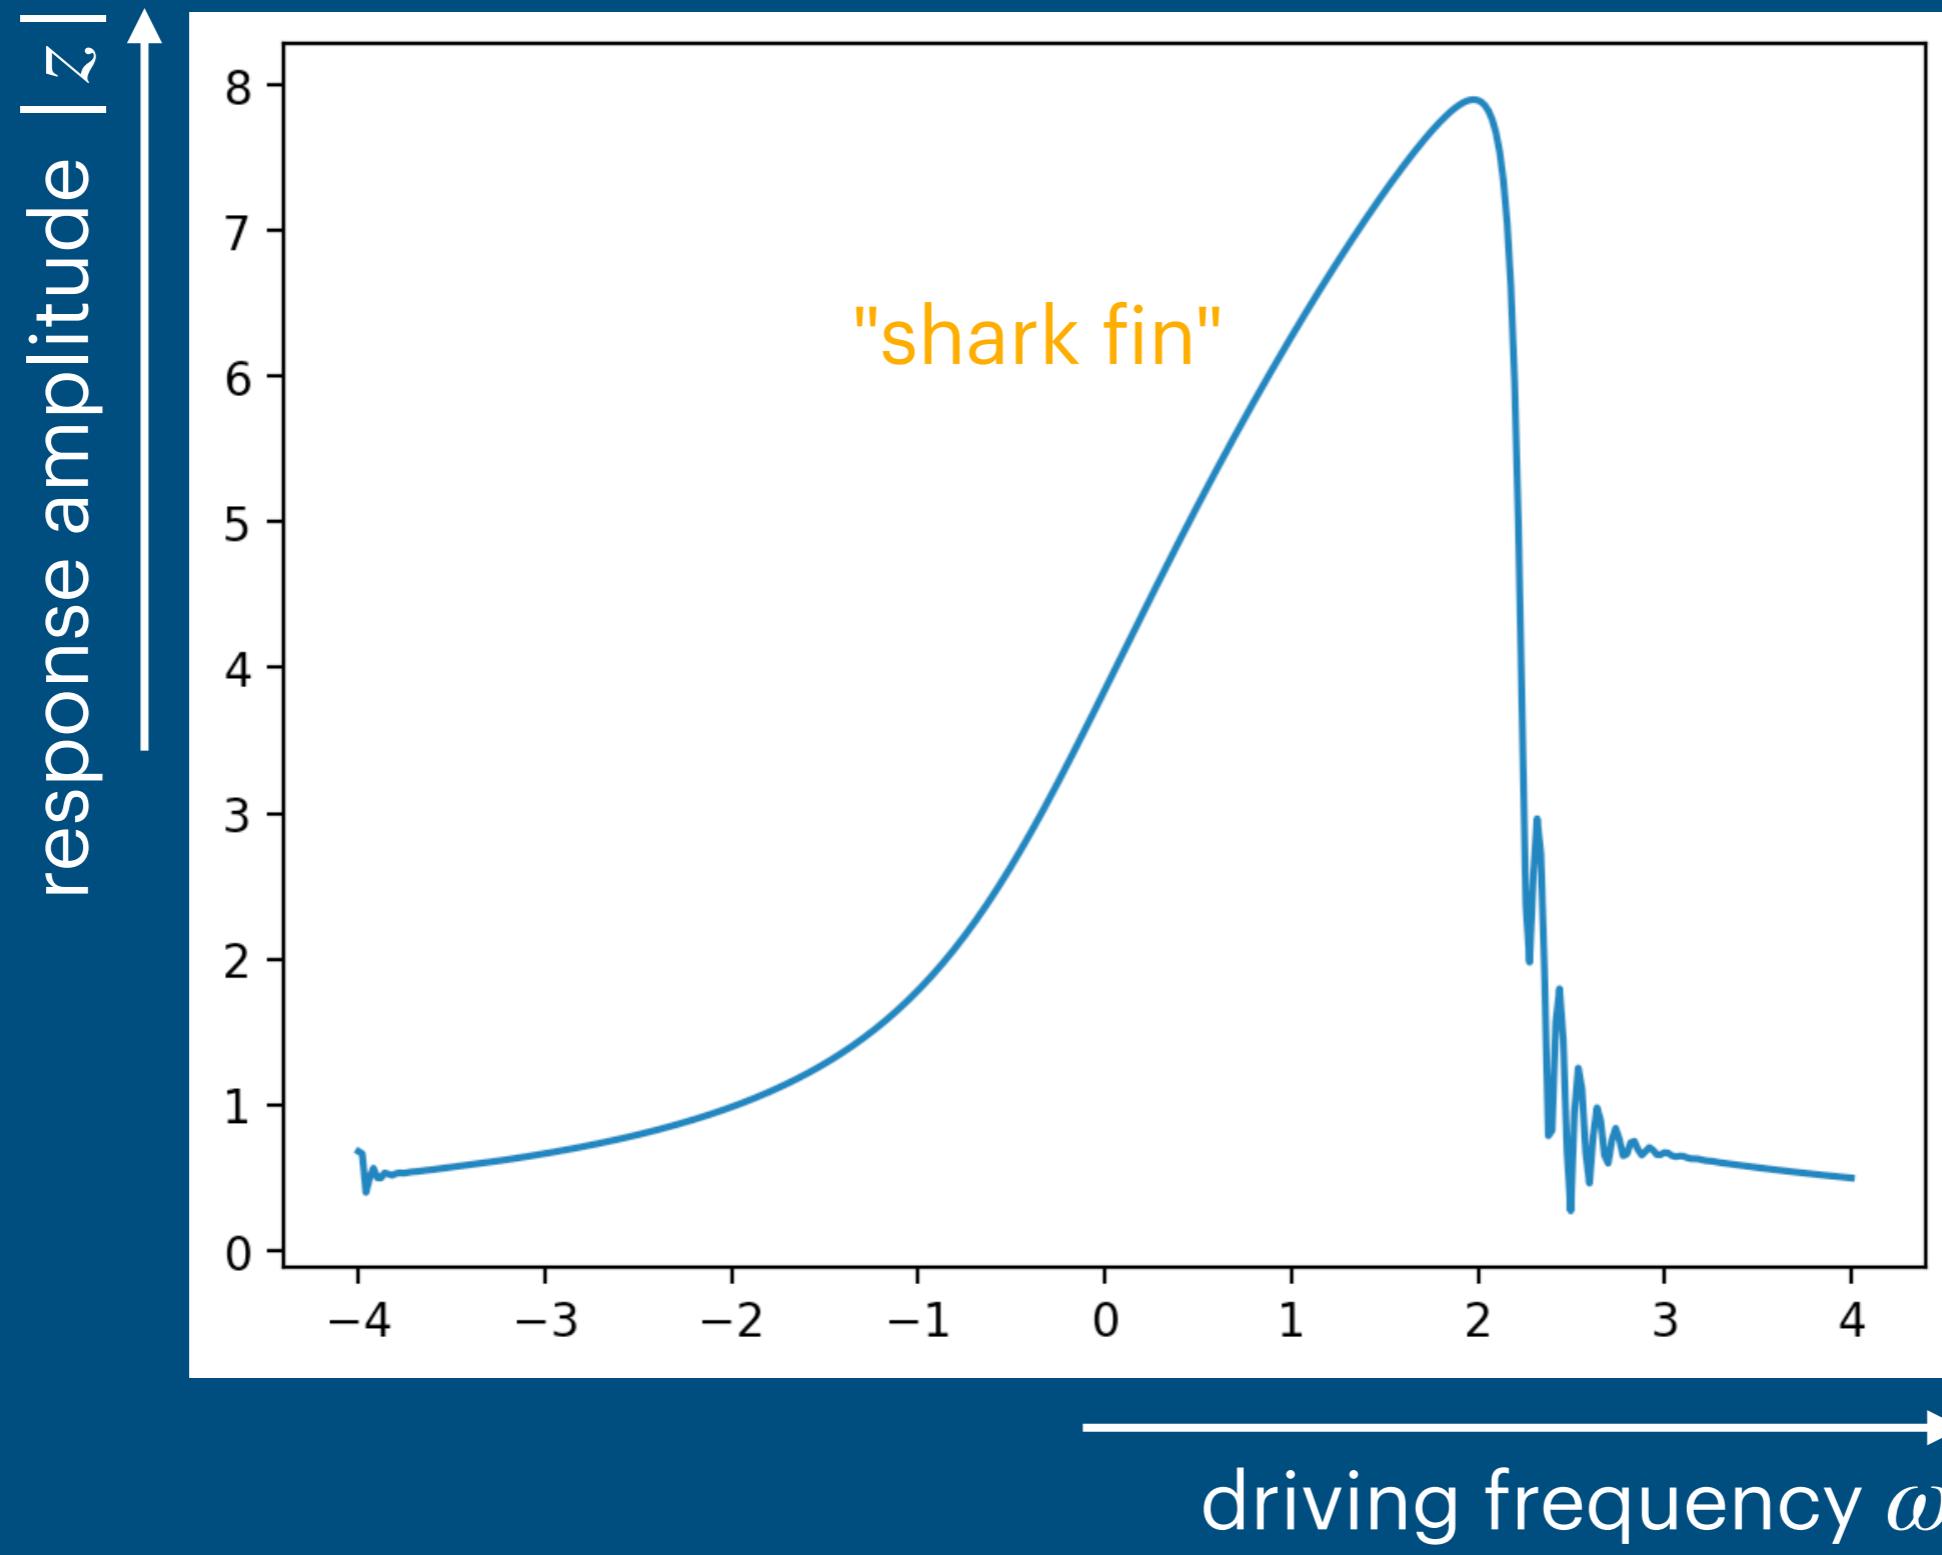
$$\dot{z} = -i(\omega_0 - \omega)z - \frac{\gamma}{2}z - i\varepsilon|z|^2z + if$$

damping      force  
 resonance freq.      nonlinearity

complex amplitude  $z = q + i\dot{q}$

# Physics example: Characterizing a physical system

## Response of a Duffing oscillator



here: drive frequency sweep in finite time (leads to wiggles)

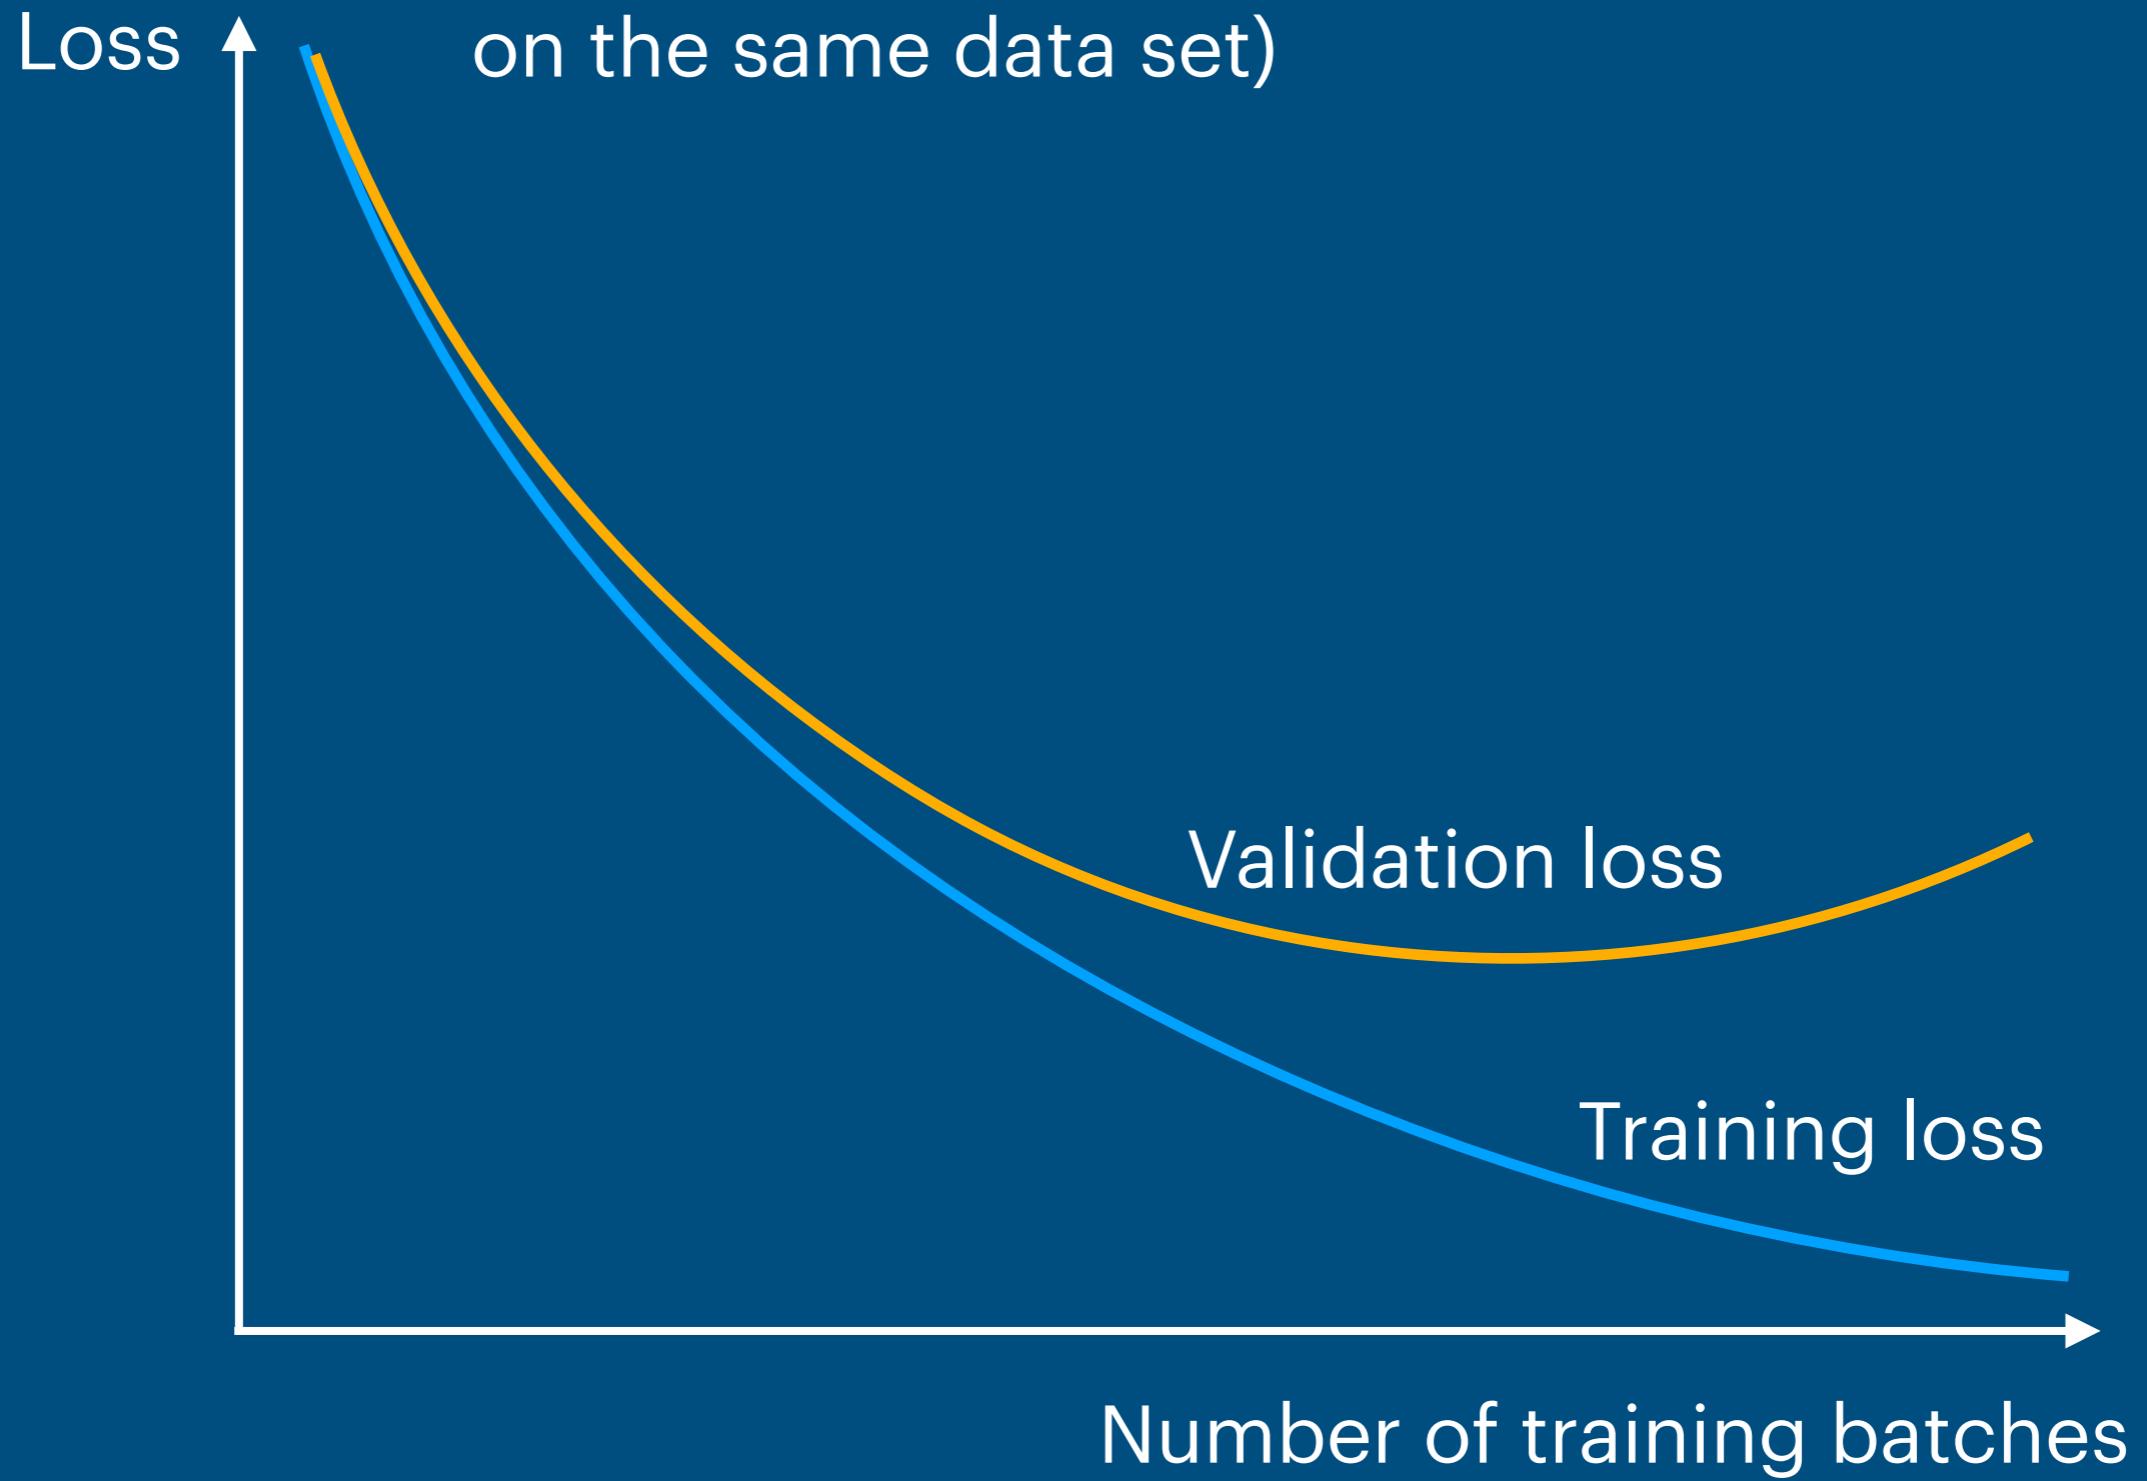
# Let us try it out!

[FlorianMarquardt.github.io/  
MachineLearningThreeEasyLessons](https://FlorianMarquardt.github.io/MachineLearningThreeEasyLessons)

(Duffing oscillator)

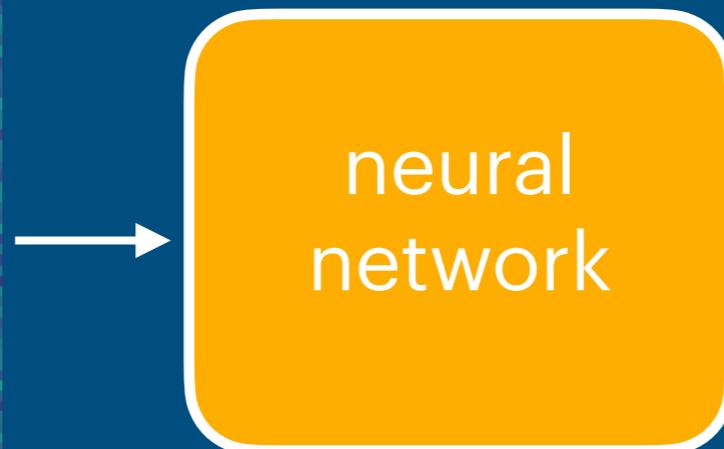
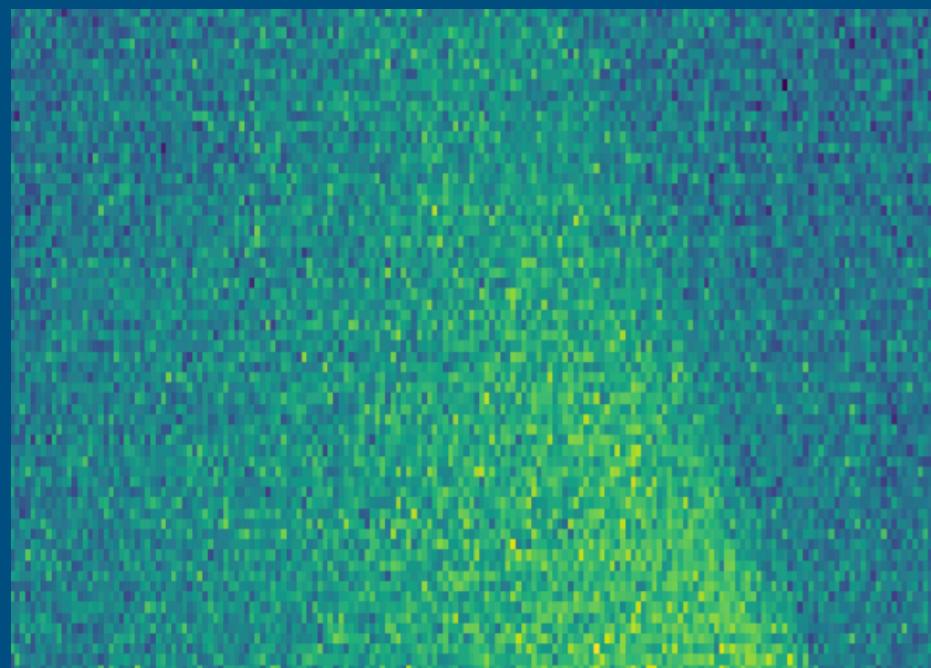
## Important Danger: Overfitting

(whenever one repeatedly trains  
on the same data set)



# Images

# Dealing with images



physical parameters:  
damping = ...  
resonance freq. = ...  
nonlinearity = ...

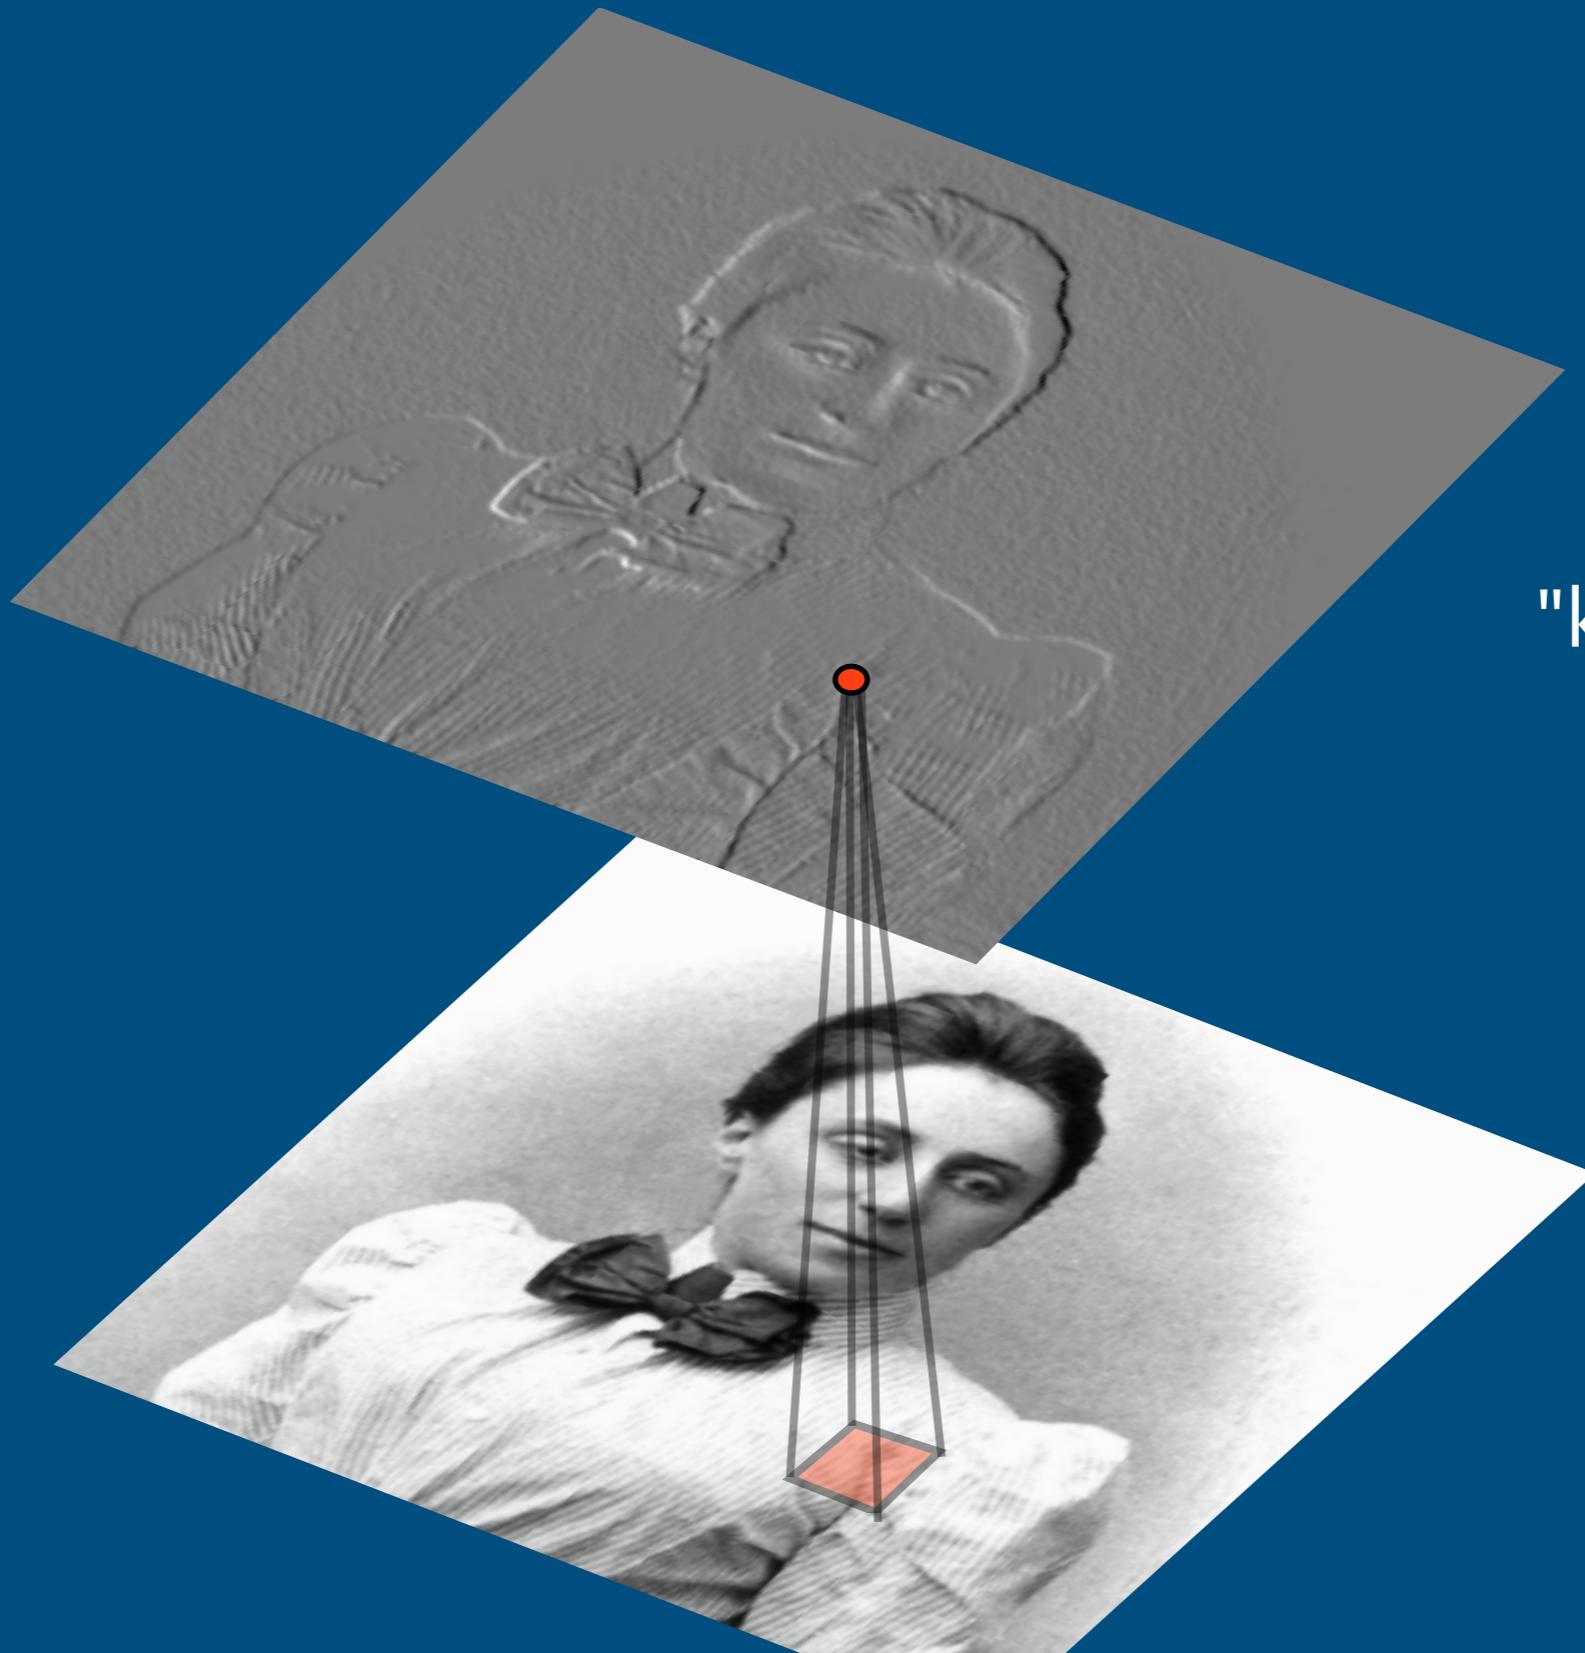
# Convolutional Neural Networks for Images



Idea: exploit translational invariance of image features to drastically reduce the number of trainable parameters and speed up training

# Convolutional Neural Networks for Images

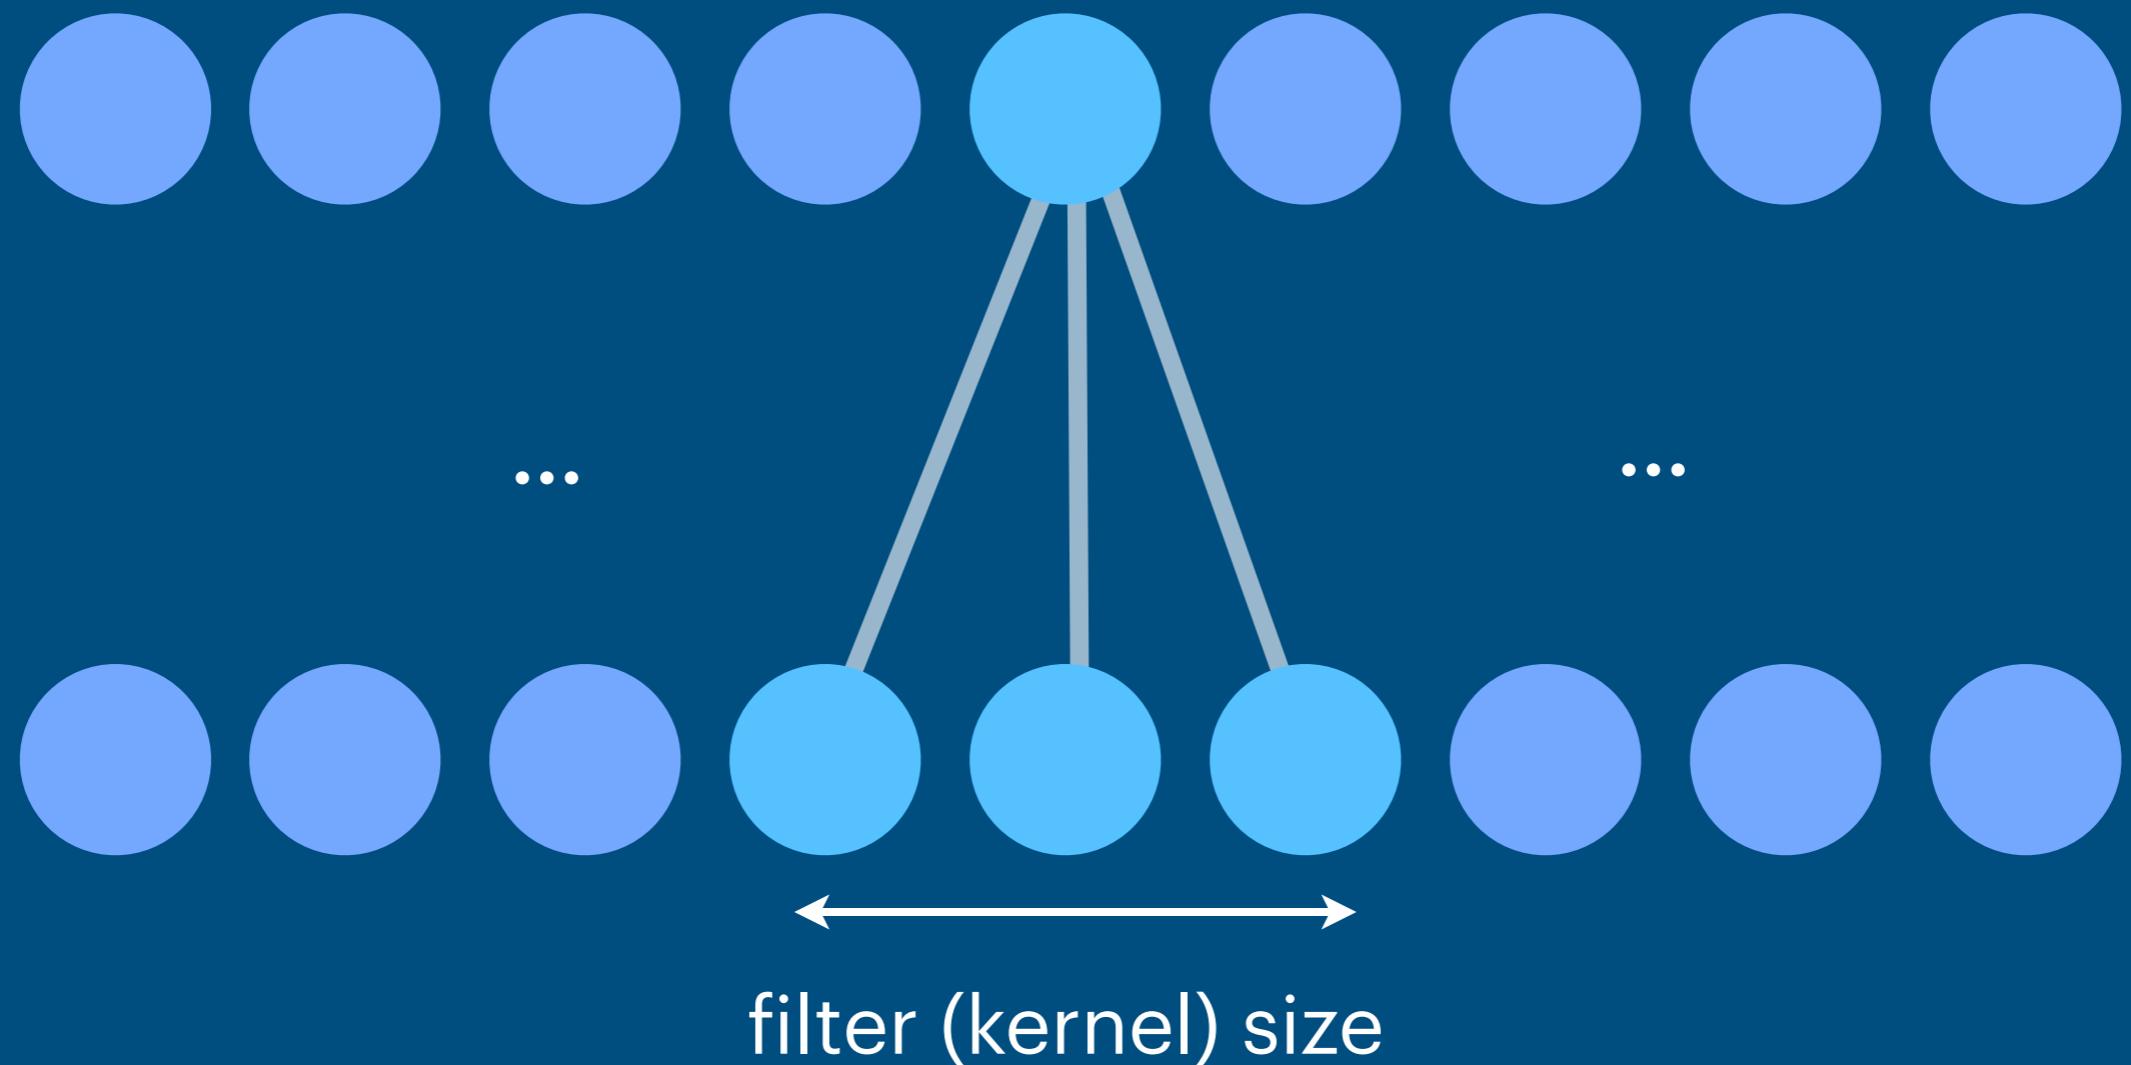
image filtering: scan some "convolution kernel"/"filter" over source image, multiply pixel values by kernel weights, add together



	+	
+		-
	-	

"kernel" (or "filter")

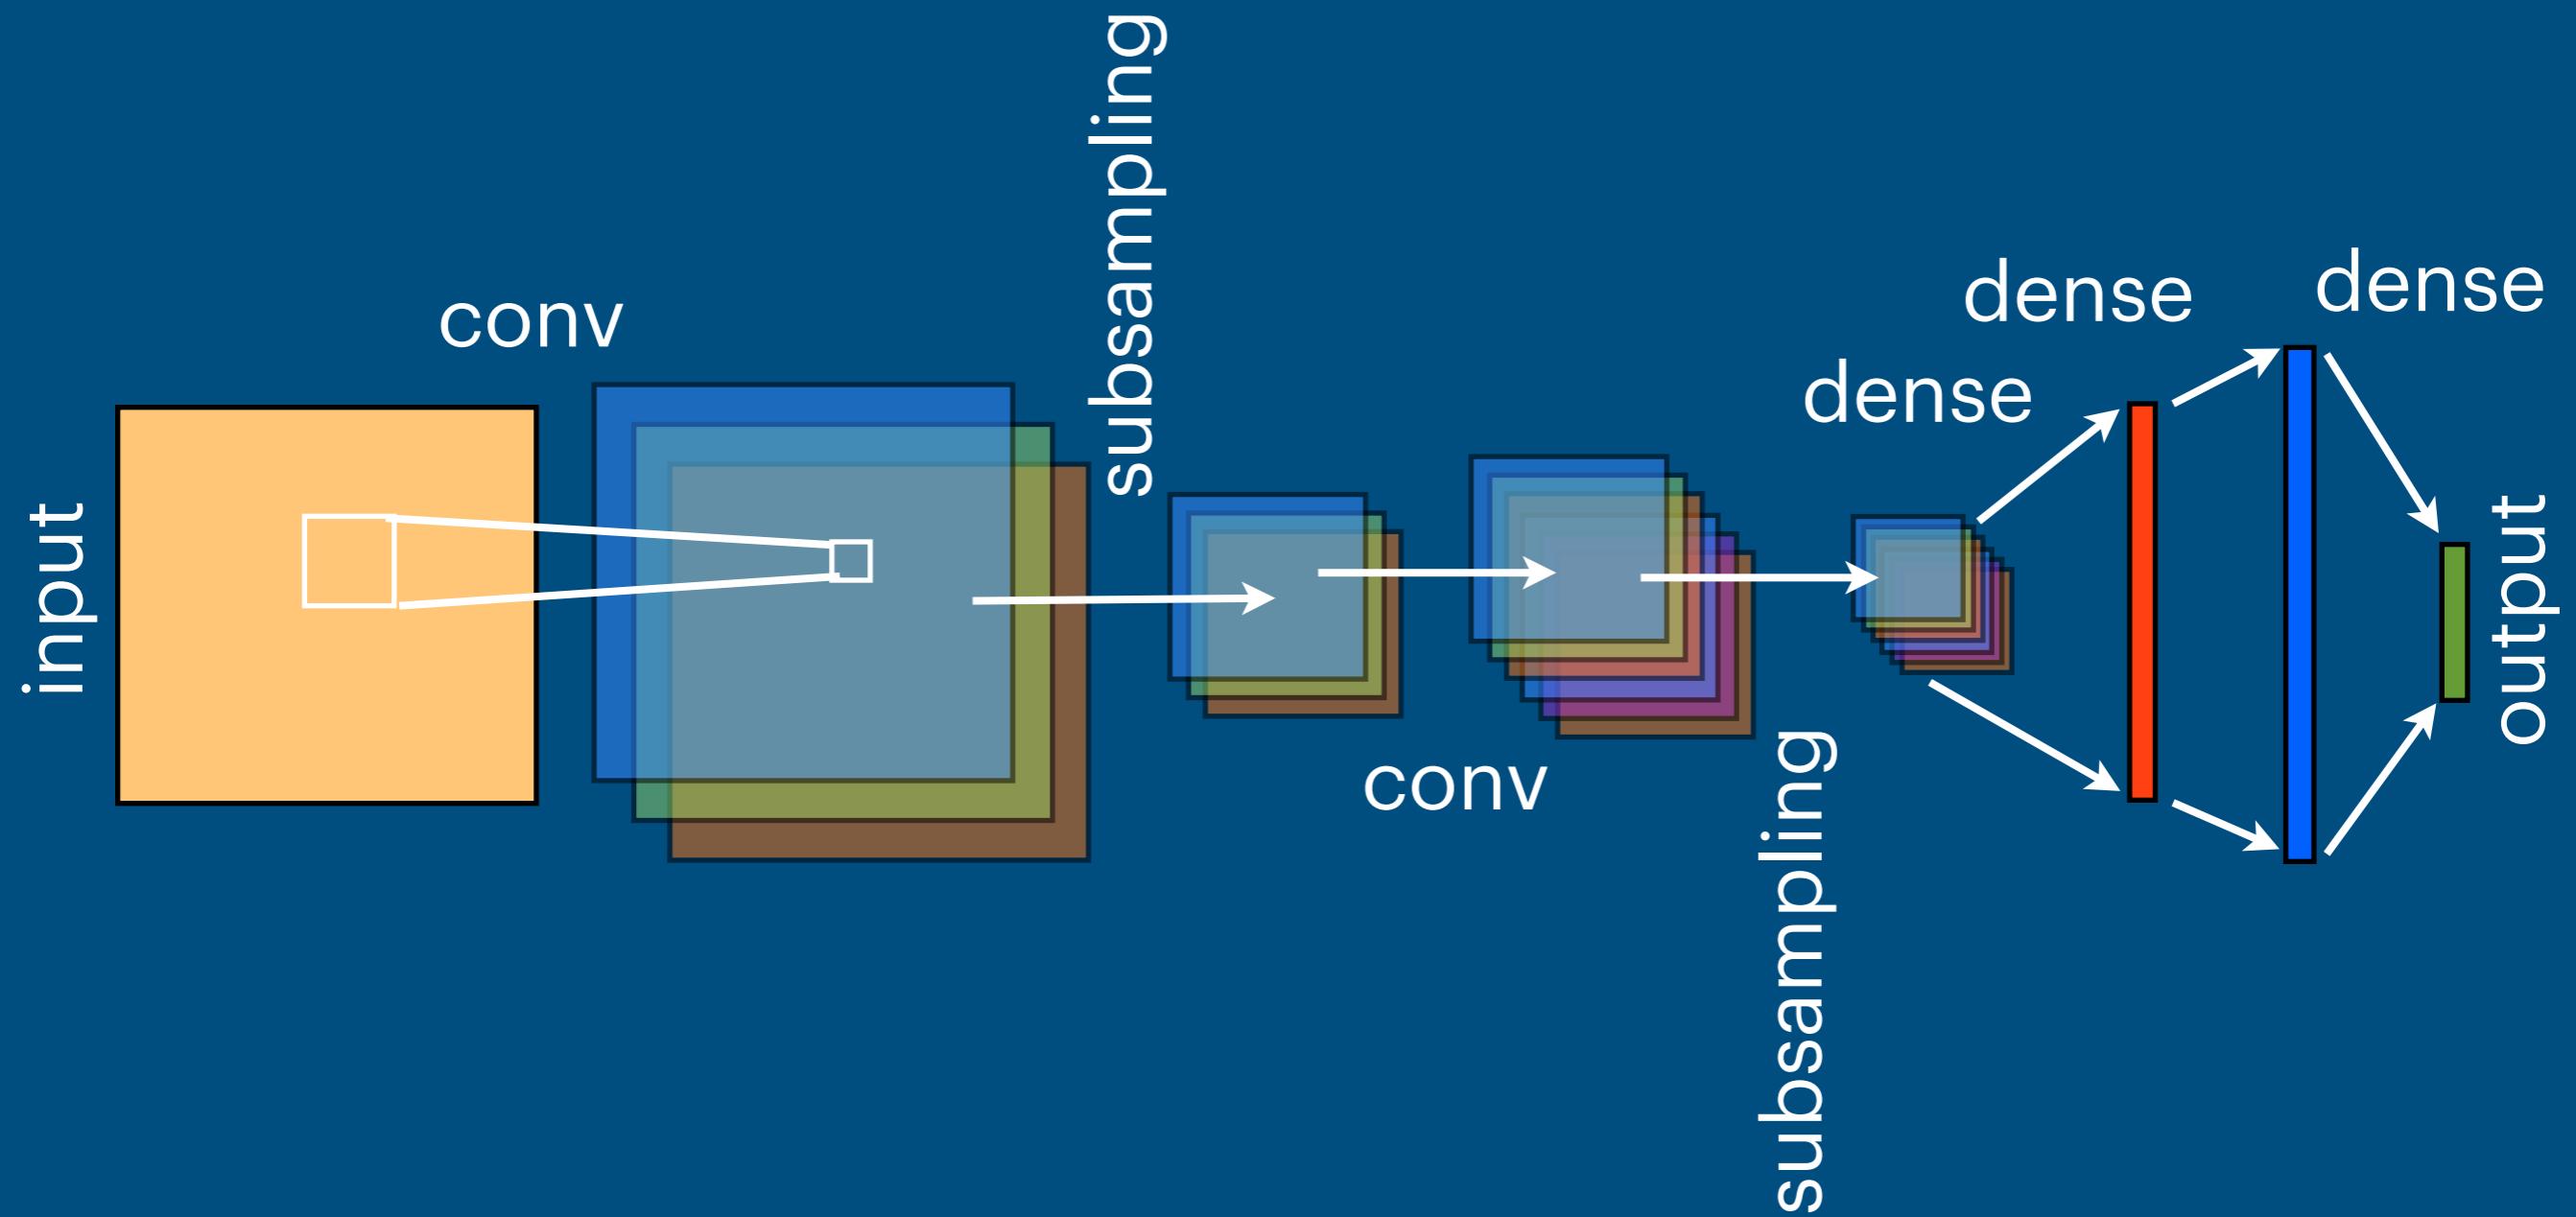
# Convolutional Neural Networks for Images



Same weights ("kernel"="filter") used for each neuron in the top layer!

Dramatically reduces number of trainable parameters

# A fully developed convolutional neural network



# Let us try it out!

[FlorianMarquardt.github.io/  
MachineLearningThreeEasyLessons](https://FlorianMarquardt.github.io/MachineLearningThreeEasyLessons)

(Lesson 3)

# Where you can go from here

Directly use what you learned to:

- speed up simulations (map initial conditions to result, sample design to properties, ...)
- interpret noisy measurement traces
- recognize experimental images
- learn resolution enhancement of images

[Technical: jax-based libraries like flax or haiku for easier construction of more advanced neural networks; or similar features in tensorflow/keras and pytorch]

## Further concepts (read about them...):

- use "autoencoders" to find compressed representations of unlabeled data
- use "residual networks" and "U-Nets" for advanced image processing
- use "recurrent networks" to analyze time series
- use "graph neural networks" for input like molecular structures and other data of variable size
- use "reinforcement learning" to discover control strategies
- use "normalizing flows" or "diffusion models" to generate new data similar to existing data
- use "transformers" to analyze and produce sequences with long dependencies

# Online lecture series by F.M., designed for physicists (florian\_marquardt\_physics YouTube channel)

## Machine Learning for Physicists (the basics)

The image shows three YouTube thumbnail cards for a lecture series. Each card has a video player interface with a timestamp at the bottom right.

- Lecture 10:** Applications in... (1:21:47) - Shows a 3D molecular model and text about protein folding.
- Lecture 9:** Boltzmann machines (1:27:59) - Shows two diagrams of energy levels and a mathematical formula for probability.
- Lecture 8:** Deep Reinforcement... (1:25:33) - Shows a flowchart of the policy gradient algorithm.

Below each thumbnail is the title, a brief description, and the view count and upload date.

Machine Learning for Physicists (Lecture 10): Applications in...	Machine Learning for Physicists (Lecture 9): Boltzmann machines	Machine Learning for Physicists (Lecture 8): Deep Reinforcement...
1.8K views • 3 years ago	2.5K views • 3 years ago	1.4K views • 3 years ago

## Advanced Machine Learning for Scientific Discovery

The image shows a YouTube channel page for "Advanced Machine Learning (lecture series)" by Florian Marquardt. The page includes a play all button and a brief description of the series.

**Advanced Machine Learning (lecture series) ► Play all**

Lecture Series 2021/22 about "Advanced Machine Learning for Physics, Science, and Artificial Scientific Discovery" by Florian Marquardt. See [https://pad.gwdg.de/...](https://pad.gwdg.de/)

The page features three large thumbnail cards for the first three lectures:

- Lecture 1: Introduction. Artificial Scientific Discovery.** (1:38:36) - Illustration of a robot and a cell.
- Lecture 2: Basic neural network structure.** (1:04:44) - Illustration of a neural network structure.
- Lecture 3: Stochastic Gradient Descent....** (1:35:15) - Illustration of a landscape with a gradient descent path.

Below each thumbnail is the title, a brief description, and the view count and upload date.

Lecture 1: Introduction. Artificial Scientific Discovery.	Lecture 2: Basic neural network structure.	Lecture 3: Stochastic Gradient Descent....
Florian Marquardt 14K views • Streamed 2 years ago	Florian Marquardt 3.6K views • Streamed 2 years ago	Florian Marquardt 2.9K views • Streamed 2 years ago